

Contour: A Practical System for Binary Transparency

Mustafa Al-Bassam
University College London
m.albassam@cs.ucl.ac.uk

Sarah Meiklejohn
University College London
s.meiklejohn@ucl.ac.uk

ABSTRACT

Transparency is crucial in security-critical applications that rely on authoritative information, as it provides a robust mechanism for holding these authorities accountable for their actions. A number of solutions have emerged in recent years that provide transparency in the setting of certificate issuance, and Bitcoin provides an example of how to enforce transparency in a financial setting. In this work we shift to a new setting, the distribution of software package binaries, and present a system for so-called “binary transparency.” Our solution, Contour, uses proactive methods for providing transparency, privacy, and availability, even in the face of persistent man-in-the-middle attacks. We also demonstrate, via benchmarks and a test deployment for the Debian software repository, that Contour is the only system for binary transparency that satisfies the efficiency and coordination requirements that would make it possible to deploy today.

1 INTRODUCTION

Historically, functional societies have relied to a large degree on trust in their governing institutions, with participants in various systems (nation states, the Internet, financial markets, etc.) trusting those in charge to follow an agreed-upon set of rules and thus provide the system with some level of integrity. In recent years, however, increasing numbers of incidents have demonstrated that integrity cannot be meaningfully achieved solely by placing trust in a small number of entities. As a result, people are now demanding more active participation in the systems with which they interact, and more accountability for the entities that govern them. The main method that has been relatively successful thus far in achieving accountability is the idea of *transparency*, in which information about the decisions within the system are made globally visible, thus enabling any participant to check for themselves whether or not the decisions comply with what they perceive to be the rules.

One of the technical settings in which the idea of transparency has been most thoroughly — and successfully — deployed is the issuance of X.509 certificates. This is partially due to the nature of these certificates (which are themselves intended to be globally visible), and partially to the many publicized failures of major certificate authorities (CAs) [17, 22]. A long line of recent research [4, 9, 19, 21, 23, 24, 28, 31] has provided and analyzed solutions that bring transparency to the issuance of both X.509 certificates (“certificate transparency”) and to the assignment of public keys to end users (“key transparency”).

Despite their differences, many of these systems share a fundamentally similar architecture [6]: after being signed by CAs, certificates are stored by *log servers* in a globally visible append-only log; i.e., in a log in which entries cannot be deleted without detection. Clients are told to not accept certificates unless they have been included in such a log, and to determine this they rely on *auditors*, who are responsible for checking inclusion of the specific

certificates seen by clients. Because auditors are often thought of software running on the client (e.g., a browser extension), they must be able to operate efficiently. Finally, in order to expose misbehavior, *monitors* (inefficiently) inspect the certificates stored in a given log to see if they satisfy the rules of the system.

To prevent clients from accepting bad certificates, such systems thus rely on monitors to expose them. Because auditors are the ones communicating with the client, however, to achieve this property an additional line of communication is needed between the auditor and monitor in the form of a *gossip* protocol [7, 27]. In such a protocol, the auditor and monitor periodically exchange information on their current and previous views of the log, which allows them to detect whether or not their views are *consistent*, and thus whether or not the log server is misbehaving by presenting “split” views of the log. If such attacks are possible, then the accountability of the system is destroyed, as a log server can present one log containing all certificates to auditors (thus convincing it that its certificates are in the log), and one log containing only “good” certificates to monitors (thus convincing them that all participants in the system are obeying the rules).

While gossiping can detect this misbehavior, it is ultimately a retroactive mechanism — i.e., it detects this behavior after an auditor has already accepted a certificate as valid and it is too late — and is thus most effective in settings where (1) no persistent man-in-the-middle (MitM) attack can occur, so the line of communication between an auditor and monitors remains open, and (2) some form of external punishment is possible, to sufficiently disincentivize misbehavior on the basis of detection. Specifically for (1), if an auditor has no means of communication that is not under an adversary’s control for the foreseeable future (a scenario we refer to as a persistent MitM attack), then the adversary may block all gossip being sent to and from the auditor, and thus monitors may never see evidence of log servers misbehaving.

Such a persistent MitM attack may be performed by an adversary who has compromised the cryptographic signing keys of the software distribution authority. This would enable them to compromise individual devices with malicious software updates, and then prevent gossiping between auditors and monitors by either using the malicious software to disable the gossiping system, or — if they control the network the device is connect to — prevent gossiping at a network level until the device stops gossiping. For example, the proposed gossip protocol for CT implements a fixed-sized pool of items to gossip, with items eventually being removed from the pool, as it would be wasteful for devices to gossip about the same information permanently [27]. An adversary would then have to carry out an attack only until this pool were emptied.

Various systems have been proposed recently that use proactive transparency mechanisms designed to operate in settings where these assumptions cannot be made, such as Collective Signing [30] (CoSi), but perhaps the most prominent example of such a system is

Bitcoin (and all cryptocurrencies based on the idea of a *blockchain*). In Bitcoin, all participants have historically played the simultaneous role of log servers (in storing all Bitcoin transactions), auditors, and monitors (in checking that no double-spending takes place). The high level of integrity achieved by this comes at great expense to the participants, both in terms of storage costs (the Bitcoin blockchain is currently over 100 GB¹) and computational costs in the form of the expensive proof-of-work mechanism required to maintain the blockchain, but several recent proposals attempt to achieve the same level of integrity in a more scalable way [20, 31]. CoSi [30] achieves this property by allowing a group of witnesses to collectively sign statements to indicate that they have been “seen,” but assumes the setup and maintenance of a Sybil-free set of witnesses, which introduces a large coordination effort.

Because of the effectiveness of these approaches, there has been interest in repurposing them to provide not only transparency for certificates or monetary transfers, but for more general classes of objects (“general transparency” [10]). One specific area that thus far has been relatively unexplored is the setting of software distribution (“binary transparency”). Bringing transparency to this setting is increasingly important, as there are an increasing number of cases in which actors target devices with malicious software signed by the authoritative keys of update servers. For example, the Flame malware, discovered in 2012, was signed by a rogue Microsoft certificate and masqueraded as a routine Microsoft software update [17]. In 2016, a US court compelled Apple to produce and sign custom firmware in order to disable security measures on a phone that the FBI wanted to unlock [13].

Challenges of binary transparency. Aside from its growing relevance, binary transparency is particularly in need of exploration because the techniques described above for both certificate transparency and Bitcoin cannot be directly translated to this setting. Whereas certificates and Bitcoin transactions are small (on the order of kilobytes), software binaries can be arbitrarily large (often on the order of gigabytes), so cannot be easily stored and replicated in a log or ledger.

Most importantly, by their very nature software packages have the ability to execute arbitrary code on a system, so malicious software packages can easily disable gossiping mechanisms, and we cannot assume that the auditor always has a means of communication that is not under an adversary’s control. Specifically, as discussed earlier a malicious adversary may perform a MitM attack to prevent gossip while presenting an auditor a malicious view of the log, and the log may itself contain a malicious software update that executes code to disable gossiping. This makes retroactive methods for detecting misbehavior uniquely poorly suited to this setting, in which clients need to know that a software package has been inspected by independent parties *before* installing it, not after. Binary transparency systems relying on such retroactive methods, based on Certificate Transparency, are currently being proposed for Firefox [1].

Our contributions. We present Contour, a solution for binary transparency that utilizes the Bitcoin blockchain to proactively

prevent clients from installing malicious software, even in the face of long-term MitM attacks. Concretely, we contribute a realistic threat model for this setting and demonstrate that Contour is able to meet it; we also show, via comparison with previous solutions, that Contour is currently the only solution able to satisfy these security properties while still maintaining efficiency and a minimal level of coordination among the various participants in the system. We also provide a prototype implementation that further demonstrates the efficiency of Contour, and finally provide an argument for its practicality via a test deployment for the Debian software repository. Putting everything together, we view Contour as a solution for binary transparency that is ready to be deployed today.

We begin in Section 4 by presenting our threat model. In addition to the goal of preventing split views, we highlight the importance of *auditor privacy*, in which auditors should not reveal the particular binaries in which they are interested (as this could reveal, for example, that a client has a susceptible version of some software), and of *availability*, in which auditors and monitors should still be able to do their job even if the original software update server loses its data or goes offline.

After then presenting the design of Contour in Section 5, we go on to analyze both its security and its efficiency in Section 6. Given the volume of related research on certificate transparency, we also present some comparisons here, and argue that ours is the first efficient solution to provide these security guarantees without requiring any coordination cost, in the form of selecting a central entity to perform authorization, or otherwise trusting some party to form a Sybil-free set of nodes.

To validate our efficiency claims, in Section 7 we describe an implementation of Contour and benchmark its performance, finding that almost all operations can be performed very quickly (on the order of microseconds), that auditors can store minimal information (on the order of kilobytes), and that arbitrary numbers of binaries can be represented by a single small (235-byte) Bitcoin transaction. We also validate our claims of real-world relevance by presenting, in Section 8, the application of Contour to the current package repository for the Debian operating system. We find that it would require minimal overhead for existing actors, and cost under 17 USD per day (even given the current high price of Bitcoin).

Finally, in Section 9 we present some possible extensions to Contour, including a discussion of how to use it to achieve general transparency, and in Section 10 we conclude.

2 RELATED WORK

There is by now a significant volume of related work on the idea of transparency, particularly in the settings of certificates, keys, and Bitcoin. We briefly describe some of this work here, and provide a more thorough comparison to the most relevant work in Section 6.3. While Contour uses similar techniques to previous solutions within these other contexts, to the best of our knowledge it is the first full deployable solution in the context of binary transparency.

In terms of certificate transparency, AKI [19] and ARPKI [4] provide a distributed infrastructure for the issuance of certificates, thus providing a way to prevent rather than just detect misbehavior. Certificate Transparency (CT) [21] focuses on the storage of certificates rather than their issuance, Ryan [28] demonstrated how

¹blockchain.info/charts/blocks-size

to handle revocation within CT, and Dowling et al. [9] provided a proof of security for it. Eskandarian et al. [11] propose how to make some aspects of gossiping in CT more privacy-friendly using zero-knowledge proofs. CONIKS [24] focuses instead on key transparency, and thus pays more attention to privacy and does not require the use of monitors (but rather has users monitor their own public keys).

In terms of solutions that avoid gossip, Fromknecht et al. [14] propose a decentralized PKI based on Bitcoin and Namecoin, and IKP [23] provides a way to issue certificates based on Ethereum. EthIKS [5] provides an Ethereum-based solution for key transparency and Catena [31] provides one based on Bitcoin. While both Catena and Contour utilize similar recent features of Bitcoin to achieve efficiency, they differ in their focus (key vs. binary transparency), and thus in the proposed threat model; e.g., Catena dismisses eclipse attacks [29] on the Bitcoin network, whereas we consider them well within the scope of a MitM attacker. Chainiac [26] is a system for proactive software update transparency based on a verifiable data structure called a skipchain. Chainiac uses a consensus mechanism based on Collective Signing (CoSi) [30], leading to the need for an authority to maintain a Sybil-free set of nodes.

Finally, in terms of more general solutions, Chase and Meiklejohn abstract CT into the general idea of a “transparency overlay” [6] and prove its security. Similarly, CoSi [20, 30] is a general consensus mechanism that shares our goal of providing transparency even in the face of MitM attacks and thus avoids gossiping, but requires setting up a distributed set of “witnesses” that is free of Sybils. This is a deployment overhead that we avoid.

3 BACKGROUND

3.1 Software distribution

Software distribution on modern desktop and mobile operating systems is managed through centralized software repositories such as the Apple App Store, the Android Play Store, or the Microsoft Store. Most Linux distributions such as Debian also have their own software repositories from which administrators can install and update software packages using command-line programs.

To reduce the trust required in these repositories, efforts such as *deterministic builds* allow users to verify that a compiled binary corresponds to the published source code of open-source software, a traditionally difficult process due to sources of non-determinism in build processes. Deterministic builds are achieved by recording the environment when building software, then replaying the behavior of this environment in later builds to achieve the same results [8]. While this prevents developers from inserting malicious code into the compiled binaries (i.e., making their code public but including a different version in the actual binary), it does not address the targeted malware threat that Contour aims to solve, in which the source code (or binary) for one targeted set of users is different from the copy received by everyone else.

3.2 Distributed ledgers

The concept of a blockchain was first used in Bitcoin, which is designed to be a globally consistent append-only ledger of financial transactions [25]. Given our limited usage of Bitcoin, we focus for brevity only on the properties that we require for Contour.

Briefly, the Bitcoin blockchain is (literally) a chain of blocks. Each block contains two components: a *header* and a list of transactions. In addition to other metadata, the header stores the hash of the block (which, in compliance with the proof-of-work consensus mechanism, must be below some threshold in order to show that a certain amount of so-called “hashing power” has been expended to form the block), the hash of the previous block (thus enabling the chain property), and the root of the Merkle tree that consists of all transactions in the block.

On the constructive side, while the scripting language used by Bitcoin is (intentionally) limited in its functionality, Bitcoin transactions can nevertheless store small amounts of arbitrary data. This makes Bitcoin potentially useful for other applications that may require the properties of its ledger, such as certifying the ownership and timestamp of a document [3]. One mechanism that allows Bitcoin to store such data is the script opcode `OP_RETURN`,² which can be used to embed up to 80 bytes of arbitrary data.

Another aspect of Bitcoin that enables additional development is the idea of an SPV (Simplified Payment Verification) client. Rather than perform the expensive verification of the digital signatures contained in Bitcoin transactions, or the checks necessary to determine whether or not double-spending has taken place, these clients check only that a given transaction has made it into some block in the blockchain. As this can be achieved using only the root hashes stored in the block headers, such clients can store only these headers (which are small) and verify only Merkle proofs of inclusion obtained from “full” nodes (which is fast), and are thus significantly more efficient than their full node counterparts.

On the destructive side, various attacks have been demonstrated that undermine the security guarantees of Bitcoin. In *eclipse* attacks [2, 16, 18], an adversary exploits the topology of the Bitcoin network to interrupt, or at least delay, the delivery of announcements of new transactions and blocks to a victim node. More expensive “51%” attacks, in which the adversary controls more than half of the collective hashing power of the network, allow the adversary to fork the blockchain, and it has been demonstrated [12] that such attacks can in fact be carried out with far less than 51% of the hashing power.

4 THREAT MODEL AND SETTING

In this section, we describe the actors in the ecosystem for software distribution transparency (Section 4.1), along with the interactions between these actors (Section 4.2), and the goals we hope to achieve in this setting (Section 4.3).

4.1 Participants

We consider a system with five types of actors: services, authorities, monitors, auditors, and clients. We describe each of these types below in the singular, but for the correct and secure functioning of a transparency overlay we require a distributed set of auditors and monitors, each acting independently.

Service: The service is responsible for producing actions, such as the issuance of a software update. In order to have these binaries authorized, they must be sent to the authority.

²en.bitcoin.it/wiki/OP_RETURN

Authority: The authority is responsible for publishing *statements* that declare it has received a given software binary from a service. These statements also claim that the authority has — in some form — published these binaries in a way that allows them to be inspected by the monitor. The authority is also responsible for placing its statements into a public *audit log*, where they can be efficiently verified by the auditor.

Monitor: The monitor is responsible for inspecting the binaries published by the authority and performing out-of-band tests to determine their validity (e.g., to ensure that software updates do not contain malware).

Auditor: The auditor is responsible for checking specific binaries against the statements made by the authority that claim they are published.

Client: The client receives software updates from either the authority or the service, along with a statement that claims the update has been published for inspection. It outsources all responsibility to the auditor, so in practice the auditor can be thought of as software that sits on the client (thus making the client and auditor the same actor, which we assume for the rest of the paper).

4.2 Interactions

In terms of the interactions between these entities, one of the main benefits of Contour — as discussed in the introduction — is that entities do not need to engage in prolonged multi-round interactions like gossiping, but rather pass messages atomically to one another. As we see in Section 6.1, this makes it significantly more expensive for an adversary to present undetected split views of a log by launching man-in-the-middle attacks. We therefore outline only non-interactive algorithms needed to generate messages, rather than interactive protocols, and wait to specify the exact inputs and outputs until we present our construction in Section 5.

Authority.commit: The authority runs this algorithm to commit statements to the audit log.

Authority.prove_incl: The authority runs this algorithm to provide a proof that a specific statement is in the audit log.

Auditor.check_incl: The auditor runs this algorithm to check the proof of inclusion for a specific statement.

Monitor.get_commits: The monitor runs this algorithm to retrieve relevant commitments from the audit log.

4.3 Goals

We break the goals of the system down into security goals (denoted with an S) and deployability goals (denoted with a D).

As discussed in the introduction, it is especially crucial in the setting of binary transparency to consider adversaries that can perform persistent man-in-the-middle attacks, as it is realistic that they would be able to compromise the client’s machine. Like certificate transparency (but unlike key transparency), we do not need to make the contents of the audit log private, as binaries are assumed to be public information, but we do need to guarantee privacy for the specific binaries that a client downloads, as this could reveal that a client has a software version susceptible to malware. Finally, even though binaries are typically large, we need to nevertheless provide a solution efficient enough to be deployed in practice.

Keeping these requirements in mind, we aim in all our security goals to defend against the specified attacks in the face of malicious authorities that, in addition to performing all the usual actions of the authority, can also perform man-in-the-middle attacks on the auditor’s network communications. If additional adversaries are considered we state them explicitly.

S1: No split views. We should prevent split-view attacks, in which the information contained in the audit log convinces the auditor that the authority published a binary, and thus it is able to be inspected by monitors, whereas in fact it is not and only appears that way in the auditor’s “split” view of the log.

S2: Availability. We should prevent attacks on availability, in which the information contained in the audit log convinces the auditor that a binary is available to be inspected by monitors, when in fact the authority has not published it or has, after the initial publication, lost it or intentionally taken it down.

S3: Auditor privacy. We should ensure that the specific binaries in which the auditor is interested are not revealed to any other parties. We thus consider how to achieve this not only in the face of malicious authorities, but in the case in which all parties aside from the auditor are malicious.

D1: Efficiency. Contour should operate as efficiently as possible, in terms of computational, storage, and communication costs. In particular, the overhead beyond the existing requirements for a software distribution system should be minimal.

D2: Minimal setup. In addition to the computational overheads, we would like as little effort — in terms of, e.g., coordination — to be done as possible in order to deploy Contour, and for it to require the minimal amount of change to the existing system.

5 DESIGN OF CONTOUR

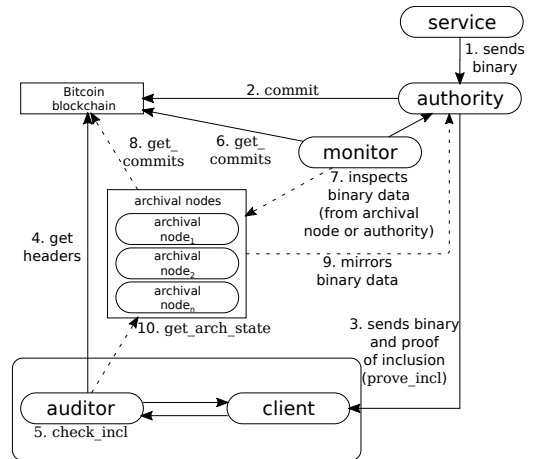


Figure 1: The overall structure of Contour. Dashed lines represent steps that are required only if archival nodes are used.

In this section we describe the overall design of Contour. An overview of all the interactions in the system can be seen in Figure 1.

5.1 Setup and instantiation

Contour and its security properties make use of a blockchain, whose primary purpose — as we see in Section 6.1 — is to provide an immutable ledger that prevents split-view attacks. Because the Bitcoin blockchain is currently the most expensive to attack, we use it here and in our security analysis in Section 6.1, but observe that any blockchain could be used in its place. An authority must initially establish a known Bitcoin address that Contour commitments are published with. As knowledge of the private key associated with the Bitcoin address is required to sign transactions to spend transaction outputs sent to the address, this acts as the root-of-trust for the authority. This address can be an embedded value in the auditor software. An initial amount of coins must be sent to the Bitcoin address to enable it to start making transactions from the address.

5.2 Logging and publishing statements

To start, the authority receives information from services; i.e., software binaries from the developers of the relevant packages (Step 1 of Figure 1). As it receives such a binary, it incorporates its hash as a leaf in a Merkle tree with root h_T . The root, coupled with the path down to the leaf representing the binary, thus proves that the authority has seen the binary, so we view the root as a batched statement attesting to the fact that the authority has seen all the binaries represented in the tree. Once the Merkle tree reaches some (dynamically chosen) threshold n in size, the authority runs the commit algorithm (Step 2 of Figure 1) as follows:

commit(h_T): Form a Bitcoin transaction in which one of the outputs embeds h_T by using OP_RETURN. One of the inputs must be a previous transaction output that can only be spent by the authority’s Bitcoin address (i.e. a standard Bitcoin transaction to the authority’s address). The other outputs are optional and may simply send the coins back to the authority’s address, according to the miner’s fees it wants to pay. (See Section 7.2 for some concrete choices.) Sign the transaction with the address’s private key and publish to the Bitcoin blockchain and return the raw transaction data, denoted tx .

Crucially, the commit algorithm stores only the root hash in the transaction, meaning its size is independent of the number of statements it represents. Furthermore, if the blockchain is append-only — i.e., if double spending is prevented — then the log represented by the commitments in the blockchain is append-only as well.

5.3 Proving inclusion

After committing a batch of binaries to the blockchain, the authority can now make these binaries accessible to clients. When a client requests a software update, the authority sends not only the relevant binary, but also an accompanying proof of inclusion, which asserts that the binary has been placed in the log and is thus accessible to monitors (Step 3 of Figure 1).

To generate this proof, the authority must first wait for its transaction to be included in the blockchain (or, for improved security, for it to be embedded k blocks into the chain). We denote the header of the block in which it was included as $head_B$. The proof then needs to convince anyone checking it of two things: (1) that the relevant binary is included in a Merkle tree produced by the authority and (2) that the transaction representing this Merkle tree

is in the blockchain. Thus, as illustrated in Figure 2, this means providing a path of hashes leading from the values retrieved from the blockchain to a hash of the statement itself.

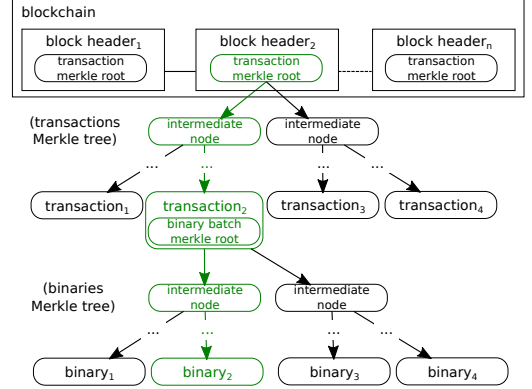


Figure 2: An example of a path of hashes leading from the block’s transactions Merkle root to the hash of bin_2 .

For a given binary bin , the algorithm `prove_incl` thus runs as follows:

prove_incl($tx, head_B, bin$): First, form a Merkle proof for the inclusion of tx in the block represented by $head_B$. This means forming a path from the root hash stored in $head_B$ to the leaf representing tx ; denote these intermediate hashes by π_{tx} . Second, form a Merkle proof for the inclusion of bin in the Merkle tree represented by tx (using the hash h_T stored in the OP_RETURN output) by forming a path from h_T to the leaf representing bin ; denote these intermediate hashes by π_{bin} . Return $(head_B, tx, \pi_{tx}, \pi_{bin})$.

5.4 Verifying inclusion

To verify this proof, the auditor must check the Merkle proofs, and must also check the authority’s version of the block header against its own knowledge of the Bitcoin blockchain. This means that the auditor must first keep up-to-date on the headers in the blockchain, which it obtains by running an SPV client (Step 4 in Figure 1). By running this client, the auditor builds up a set $S = \{head_{B_i}\}_i$ of block headers, which it can check against the values in the proof of inclusion. This means that, for a binary bin , `check_incl` (Step 5 in Figure 1) runs as follows:

check_incl($S, bin, (head_B, tx, \pi_{tx}, \pi_{bin})$): First, check that $head_B \in S$; output 0 if not. Next, extract h_T from tx (using the hash stored in the OP_RETURN output), form $h_{bin} \leftarrow H(bin)$, and check that π_{bin} forms a path from the leaf h_{bin} to the root h_T . Finally, form $h_{tx} \leftarrow H(tx)$, and check that π_{tx} forms a path from the leaf h_{tx} to the root hash in $head_B$. If both these checks pass then output 1; otherwise output 0.

As well as verifying the inclusion proof, the auditor must also check that the address that the proof’s transaction was sent from matches the authority’s address (i.e. one of the transaction inputs must be a previous transaction output that can only be spent by the authority’s address).

5.5 Ensuring availability

Independently of auditors, monitors must retrieve all commitments associated with the authority from the blockchain and mirror their binaries (Steps 6 and 7 of Figure 1). This means `get_commits` runs as follows:

`get_commits()`: Retrieve all transactions in the blockchain sent with the authority’s address, and return the hashes stored in the `OP_RETURN` outputs.

After checking the binaries against their commitments, the monitors then inspect them — to, e.g., ensure they are not malware — in ways we consider outside of the scope of this paper.

While the system we have described thus far functions correctly and allows monitors to detect if an authority has committed to a binary but not published it, in order to make the binaries themselves available for inspection, we assume the monitors can mirror the authority’s logs. It therefore fails to satisfy our goal of availability in the event that the authority goes down at some point in time.

We thus consider the case where the authority commits binaries to the blockchain, but — either intentionally or because it loses the data sometime in the future — does not supply the data to monitors. While this is detectable, as monitors can see that there are commitments in the blockchain with no data behind them, to disincentive this behavior requires some retroactive real-world method of punishment. More importantly, it prevents the monitor from pinpointing specific bad actions, such as malicious binaries, and thus from identifying potential victims of the authority’s misbehavior.

Because of this, it is thus desirable to not only enable the detection of this form of misbehavior, but in fact to prevent it from happening in the first place. One way to achieve this is to have auditors mirror the binary themselves and send it to monitors before accepting it, to ensure that they have seen it and believe it to be benign. While this would be effective, and is arguably practical in a setting such as Certificate Transparency (modulo concerns about privacy) where the objects being sent are relatively small, in the setting of software distribution — where the objects being sent are large binaries — it is too inefficient to be considered.

Instead, we propose a new actor in the ecosystem presented in Section 4: archival nodes, or *archivists*, that are responsible for mirroring all data from the authority (Steps 8 and 9 in Figure 1). To gain the extra guarantee that the data is available to monitors, auditors may thus use any archival nodes of which they are aware to check their state (i.e., the most recent block header for which they have data from the authority) and ensure that they cover the block headers relevant to the proofs they are checking (Step 10 in Figure 1). This means adding the following two algorithms to the list in Section 4.2:

`Archivist.get_commits()`: The archivist runs this algorithm to access the commitments made by the authority, just as is done by the monitor (using the same algorithm).

`Auditor.get_arch_state()`: The auditor (optionally) runs this algorithm to obtain the state of any archivists of which it is aware. This is simply the latest block header for which the archival node has mirrored the data behind the commitments held within.

Using archival nodes makes it possible to continue to pinpoint specific bad actions in the past (e.g., the publication of malware),

even if the authority loses or stops providing this data, but we stress that their usage is optional and affects only availability. Essentially, archival nodes allow for a more granular detection of the misbehavior of an authority, but do come at the cost of requiring additional nodes to store a potentially large amount of data. If such granularity is not necessary, or if the system has no natural candidates with the necessary storage requirements, then archival nodes do not need to be used and the system still remains secure. In Section 8 we explore the role of the archival nodes in the Debian ecosystem and discover that, while the storage costs are indeed expensive, there is already at least one entity playing this role.

6 EVALUATION

In this section, we evaluate Contour in terms of how well it meets the security goals (Section 6.1) and deployability goals (Section 6.2) specified in our threat model in Section 4.3. We also compare it with respect to previous solutions in Section 6.3, and argue that it is the only system to achieve all our goals.

6.1 Security goals

6.1.1 No split views (S1). In order to prevent split views, we rely on the security of the Bitcoin blockchain and its associated proof-of-work-based consensus mechanism. If every party has the same view of the blockchain, then split views of the log are impossible, as there is a unique commitment to the state of the log at any given point in time. The ability to prevent split views therefore reduces to the ability to carry out attacks on the Bitcoin blockchain.

If, for whatever reason, the adversary cannot carry out an eclipse attack, then it can perform a split-view attack only if it can fork the Bitcoin blockchain. This naïvely requires it to control 51% of the network’s mining power, which we estimate would cost roughly 2043M USD in electricity and hardware costs as of December 2017 (see Appendix A for the analysis). Regardless of the exact number, it is generally agreed that carrying out such an attack is prohibitively expensive.

If an eclipse attack is possible, due to the adversary’s MitM capability, the adversary can “pause” the auditor at a block height representing some previous state of the log, and can prevent the auditor from hearing about new blocks past this height. It is then free to mine blocks at its own pace, and so performing a split-view attack would be significantly cheaper. As a key distinguishing property of Contour’s threat model is that split-view attacks should be prevented even in the face of an adversary that can carry out such attacks, it is important to consider the nuances and costs of this attack, especially as we are not aware of any previous literature considering the costs of eclipse attacks on Bitcoin nodes.

The cost of performing an eclipse attack depends on how much time the adversary has to perform a split-view attack, as the hash rate depends on the number of mining rigs available. As a rough estimate (see Appendix A for calculations), if auditors consider a Bitcoin transaction to be confirmed after 6 blocks (the standard for most Bitcoin wallets), then as of December 2017 the attack would cost 8.3M USD if the adversary wants to perform the attack within a week. This would mean, however, that the auditor would receive a new block only every 1.4 days, which would be detectable as an eclipse attack. If auditors conservatively require that new blocks

arrive in intervals of up to three hours before assuming that they are the victim of an eclipse attack, then as of December 2017 an attack would cost roughly 91.8M USD.

6.1.2 Availability (S2). While the decentralized (and thus fully replicated) nature of the blockchain can guarantee availability, it guarantees these properties only with respect to the commitments to statements made by the authority, rather than with respect to the statements — and thus the binaries — themselves. As discussed in Section 5.5, the use of the blockchain thus does not guarantee that binaries are actually available for inspection, or will continue to be into the future.

Even just using monitors, Contour can already detect that an authority committed a statement without making the statement data (i.e., the actual binaries) available. Using the archival nodes introduced in Section 5.5, we can achieve a stronger notion of availability — in which as long as the binaries have been published at some point they can be retrieved indefinitely into the future — as long as these nodes are honest about whether or not they have mirrored the relevant data.

In binary transparency, many ISPs and hosting providers already provide their customers local mirrors of Debian repositories. We therefore envision that ISPs can act as archival nodes on behalf of their hosting clients, which creates a decentralized network of archival nodes. We elaborate on the overheads required to do so in Sections 7.2 and 8.

6.1.3 Auditor privacy (S3). Recall from Section 4.2 that one of the goals of Contour was to avoid prolonged interactions and engage only in the atomic exchange of messages. In particular, the auditor receives pre-formed proofs of inclusion from the authority (as opposed to having to request them for specific binaries, as they would in all certificate and key transparency systems), retrieves commitments directly from the blockchain, does not engage in any form of gossip with monitors, and receives the latest block hash from archival nodes without providing any input of its own. We thus achieve privacy by design, as at no point does the auditor reveal the statements in which it is interested to any other party.

One particular point to highlight is that Contour achieves auditor privacy despite the fact that auditors run SPV clients, which are known to potentially introduce privacy issues due to the use of Bloom filtering and the reliance on full nodes. This is because the proofs of inclusion contain both the raw transaction data and the block header, so the auditor does not need query a full node for the inclusion of the transaction and can instead verify it itself (and, as a bonus, saves the bandwidth costs of doing so).

6.2 Deployability goals

6.2.1 Efficiency (D1). Table 1 summarizes the computational complexity of each of the operations required to run Contour, and Table 2 summarizes the size complexity (which in turn informs the bandwidth requirements, as we explore further in Section 7.2).

As we will see in Sections 7.2 and 8, in real deployments of Contour there are already significant storage costs for the authority and archival nodes, as they must store the full set of binaries. It therefore does not impose a significant additional burden to have them perform relatively inefficient (i.e., linear in n) operations or

Operation	Time complexity
commit	$O(n_S)$
prove_incl (one-time)	$O(\log(n_T))$
prove_incl (per statement)	$O(\log(n_S))$
check_incl	$O(\log(n_S) + \log(n_T))$

Table 1: Asymptotic computational costs for the operations of Contour, where n_S is the number of statements in a batch and n_T is the number of transactions in a block.

Object	Size Complexity
Inclusion proof	$O(\log(n_S) + \log(n_T))$
Log commitment (tx)	$O(1)$
Archival node data overhead	$O(n_S)$

Table 2: Asymptotic storage costs for the objects in Contour, where n_S is the number of statements in a batch and n_T is the number of transactions in a block.

store relatively inefficient objects. As for the end-user devices on which the auditor is run, we impose a relatively minimal performance overhead (with everything logarithmic in n_S and/or n_T), and confirm this in Section 7.2.3.

6.2.2 Minimal setup (D2). In terms of coordination, the only setup requirement in Contour is the role of the archival nodes, as the rest is just a matter of adding software. As we will see in Section 8 when we look at Debian, in some settings there are already natural candidates for these actors, but if these actors are not interested in the guarantees of Contour then we can still deploy it without requiring the existing actors to change their behavior. More importantly, there are no trust requirements placed on these nodes to prevent log equivocation: even if archival nodes misbehave, monitors can still individually detect misbehavior by an authority that publishes commitments but not the underlying data. This is in stark contrast to previous solutions that require the initial establishment of a semi-trusted set of nodes.

6.3 Comparison with existing solutions

To fully pinpoint both the benefits and tradeoffs of Contour, we compare it with several known systems designed to provide transparency. In particular, we consider the tradeoffs as compared to Certificate Transparency (CT), Collective Signing (CoSi) [30], CONIKS [24], and Bitcoin. We summarize these tradeoffs in Table 3.

Looking at Table 3, we first mention that the efficiency numbers for CoSi are somewhat misleading, as there is no global log and thus no notion of checking inclusion in the log; this is why we list the efficiency costs as constant. In fact, only Bitcoin and Contour ensure a globally consistent ledger, as certificates are stored in a distributed set of logs in CT and CONIKS and there is no proposed method for achieving consensus amongst them.

Arguably the main benefit of both CT and CONIKS is their efficiency, as the auditor is required to store only a single hash. The tradeoff, however, is that they cannot prevent the authority from launching a split-view attack, but instead rely on gossiping mechanisms to detect such misbehavior after the fact. As discussed in

	Security goals (S1-S3)			Deployability goals (D1-D2)		
	Split views	Availability	Auditor privacy	Efficiency (cost)	Efficiency (size)	Minimal setup
CT	detect	no*	no	$\log(n)$	1	no
CoSi	prevent	yes*	yes	1	1	no
CONIKS	detect	no	no	$\log(n)$	1	no
Bitcoin	prevent	yes	yes	n	n	yes
Contour	prevent	yes	yes	$\log(n)$	b	yes

Table 3: A comparison between existing solutions and Contour in terms of the five goals presented in Section 4.3. For efficiency, we measure the asymptotic costs for the auditor in terms of both the computations it must perform (‘cost’) and the data it must store (‘size’). We use n to denote the number of statements and b to denote the number (but not size) of blocks in the Bitcoin blockchain. For CoSi, availability is not a explicit requirement, but can be satisfied as long as at least one witness retains the data, and for CT it is not satisfied by the basic design but could be if auditors and monitors gossiped about certificates.

the introduction, this is problematic in a setting — like binary transparency — in which adversaries can launch persistent man-in-the-middle attacks. These systems also do not achieve robust privacy for the auditor, as it must periodically reveal information to the authority (or the monitor) about the objects in which it is interested.

The other main tradeoff we observe is, perhaps unsurprisingly, between efficiency and setup costs. The first three systems all require the establishment of some initial set of distributed entities — in the case of CT, log servers are essentially authorized by Google, in the case of CONIKS, identity providers are chosen by users and listed in a PKI, and in the case of CoSi, witnesses must form a Sybil-free set — that are trusted to some extent (if not individually, then as a group). We require no such setup, which means Contour is much more easily integrated into existing systems.

In contrast, in both Bitcoin and Contour, the blockchain is maintained by a decentralized network and is not subject to intervention by central authorities. While Contour mitigates the inefficiency of Bitcoin, it still requires the auditor to store some information from all the block headers. We show in the next two sections that Contour is nevertheless efficient enough to be practical, but leave it as an interesting open problem to investigate to which extent these tradeoffs between efficiency and decentralization are inherent.

7 IMPLEMENTATION AND PERFORMANCE

To test Contour and analyze its performance, we have implemented and provided benchmarks for a prototype Python module and toolset that developers can use. We have released the implementation as an open-source project.³

7.1 Implementation details

The implementation consists of roughly 1000 lines of Python code, and provides a set of developer APIs and corresponding command-line tools. We used SHA-256 as the hashing algorithm to build Merkle trees, and modified versions (for Bitcoin compatibility) of an existing Merkle tree implementation (<https://github.com/jvsteiner/merkletree>) and a Python-based Bitcoin library `pycoinnet` (<https://github.com/richardkiss/pycoinnet/>) in order to develop our Merkle tree and SPV client, respectively.

Authority: We provide API calls for `Authority.commit`, which commits batches of statements to the Bitcoin blockchain, and

`Authority.prove_incl`, which allows it to generate inclusion proofs for individual statements.

Auditor: We provide an API call for `Auditor.check_incl`, which allows end-user software to verify proofs of inclusion. We also provide an `Auditor.sync` call that uses the Bitcoin SPV protocol to download and verify all the block hashes in the Bitcoin blockchain, so that inclusion proofs can be efficiently verified independently of third parties. (This call needs to be run only once.)

Monitor: We provide an API call for `Monitor.get_commits`, which gets all the statement batches associated with a specific authority. Monitors can then use these commitments to check the validity of the statement data (which they can retrieve from the authority or an archival node using a web server), and do whatever manual inspection is necessary; we consider this functionality outside of the scope of this paper.

Archival node: The archival node API can be used to operate an archival node, by specifying the authority’s Bitcoin address and web address where statement data is published. The archival state and mirrored statement data is stored as flat files on disk, allowing the archival node to provide access to auditors and monitors by running a web server. By accessing the archival state via a HTTPS server, auditors can securely authenticate the state of the archival node using public-key cryptography.

7.2 Performance

To evaluate the performance of our implementation, we tested all the operations listed above on a laptop with an Intel Core i5 2.60 GHz CPU and 12 GB of RAM, that was connected to a WiFi network with an Internet connection of 5 Mbit/s. We also assume that a batch to be committed contains 1 million statements, although as was seen in Table 1 — and will be confirmed later on in Figure 3 — these numbers scale as expected (either logarithmically or linearly), so it is easy to extrapolate the results for other batch sizes given the ones we present here.

We consider the complexity of these operations in terms of their computational, storage, and bandwidth requirements. A summary of our timing benchmarks can be found in Table 4, and our bandwidth requirements are in Table 5.

7.2.1 Number of transactions per block. The overhead of both generating and verifying a proof of inclusion is dependent on the number of transactions in a Bitcoin block. To capture the worst-case

³<https://github.com/musalbas/contour>

Operation	Time (μ s)	σ (μ s)
commit	5.93 (s)	0.297 (s)
prove_incl (one-time)	8.5	5.4
prove_incl (per statement)	12	6.4
check_incl	224	62.14

Table 4: Average time of individual operations, and standard deviation σ , when the batch size is 1M. The timings for commit were averaged over 20 runs, and for prove_incl and check_incl over 1M runs. The timings for commit are in bold to emphasize that they are in seconds, not microseconds.

Operation	Bandwidth
Authority.commit (using APIs)	1 MB
Authority.commit (one-time setup for full node)	126 GB
Authority.commit (using full node)	235 B
Auditor.sync	37.4 MB
Auditor.prove_incl	1.3 kB

Table 5: The bandwidth cost of operations, when the batch size is 1M. The cost of Authority.commit depends on whether or not the authority is running a full Bitcoin node or relying on third party APIs. For running a full node, there is a one-time setup cost to synchronize the blockchain.

scenario, we consider the maximum number of transactions that can fit into a block. Currently, the Bitcoin block size limit is 1 MB, up to 97 bytes of which is non-transaction data.⁴ The minimum transaction size is 166 bytes,⁵ so the upper bound on the number of transactions in a given block is 6,023. While this is far higher than the number of transactions that Bitcoin blocks currently contain,⁶ we nevertheless use it as a worst-case cost and an acknowledgment that Bitcoin is evolving and blocks may grow in the future.

7.2.2 Authority overheads. To run commit and prove_incl, an authority must have access to the full blocks in the Bitcoin blockchain, as well as the ability to broadcast transactions to the network. Rather than achieve these by running the authority as a full node, our implementation uses external blockchain APIs supplied by blockchain.info and blockcypher.com. This decision was based on the improved efficiency and ease of development for prototyping, but it does not affect the security of the system: authorities do not need to validate the blockchain, as invalid blocks from a dishonest external API simply result in invalid inclusion proofs that are rejected by the auditor.

To run commit, an authority must first build the Merkle tree containing its statements. Sampled over 20 runs, the average time to build a Merkle tree for 1M statements was 5.9 s ($\sigma = 0.29$ s). After building the tree, an authority next embeds its root hash (which is 32 bytes) into an OP_RETURN Bitcoin transaction to broadcast to the network. Sampled over 1,000 runs, the average time to generate this transaction — in the standard case of one input and two outputs, one for OP_RETURN and one for the authority’s change — was 0.03 s ($\sigma = 0.007$ s). The average total time to run commit was thus 5.93 s,

as seen in Table 4, and it resulted in 235 bytes (the size of the transaction) being broadcast to the network.

Next, to run prove_incl, the authority proceeds in two phases: first constructing the Merkle proof for its transaction within the block where it eventually appears, and next constructing the Merkle proof for each statement represented in a transaction. The time for the first phase, averaged over 1M runs and for a block with 6,023 transactions (our upper bound from Section 7.2.1), was 8.5 μ s. This is denoted “one-time” in Table 4 as it is done only once per batch. The time for the second phase, averaged over 1M runs, was 12 μ s for each individual statement (thus denoted “per statement” in Table 4). Generating inclusion proofs for all the statements in the batch would thus take around 12 s. In terms of bandwidth and storage, a block up to 1 MB in size needs to be downloaded in order to generate the inclusion proof from the block’s transaction Merkle tree. In terms of the memory costs, the size of the Merkle tree for 1M leaves in memory is 649 MB.

Additionally, in order to ensure that its transaction makes it into a block quickly, the authority may want to pay a fee. The recommended rate as of December 5 2017 is 154 satoshis/byte (<https://bitcoinfees.info>), so for a 235-byte transaction the authority can expect to pay 36,190 satoshis. As of December 5 2017, this is roughly 4.21 USD. We stress, however, that the Bitcoin price is notoriously volatile (for example, the same transaction would have cost only 0.28 USD at the beginning of 2017), so this and all other costs stated in fiat currency should be taken with a grain of salt.

7.2.3 Auditor overheads. For the auditor, we considered two costs: the initial cost to retrieve the necessary header data (sync), and the cost to verify an inclusion proof (check_incl). We do not provide benchmarks for the Auditor.get_arch_state call, as this is a simple web request that returns a single 32-byte hash.

To run sync, auditors use the Bitcoin SPV protocol to download and verify the headers of each block, which are 80 bytes each. As of December 5 2017, there are 497,723 valid mined blocks, which equates to 39.8 MB of block headers. Once downloaded, however, the auditor needs to keep only the 32-byte block hash, so only 15.9 MB of data needs to be stored on disk. Going forward, the Bitcoin network generates approximately 144 blocks per day, so the amount of downloaded data will be 11.5 kB daily, and the amount of stored data will increase by 4.6 kB daily.

To verify the validity of the block headers in the chain, the client must perform one SHA-256 hash per block header; averaged over five runs, it took us 116 seconds for the Python SPV client to download and verify all the block headers. This initial bootstrapping process needs to be performed only once per auditor.

To run check_incl, we again use our upper bound from Section 7.2.1 and assume every block contains 6,023 transactions. This means the inclusion proof contains: (1) an 80-byte block header; (2) the raw transaction data, which is 235 bytes; (3) a Merkle proof for the transaction, which consists of $\log(6023) - 1$ 32-byte hashes (the root hash is already provided in the block header); and (4) a Merkle proof for the statement, which consists of $\log(1000000) - 2$ 32-byte hashes (the root hash is already provided in the transaction data, and the auditor computes the statement hash itself). The total bandwidth cost is therefore around 1275 bytes. Averaged over 1M

⁴<https://en.bitcoin.it/wiki/Block>

⁵https://en.bitcoin.it/wiki/Maximum_transaction_rate

⁶<https://blockchain.info/charts/n-transactions-per-block>




Figure 3: The time to verify an inclusion proof with varying batch sizes, averaged over 100K runs.

runs, the time for the auditor to verify the inclusion proof was $224 \mu\text{s}$ ($\sigma = 62.14 \mu\text{s}$).

To confirm that the time to run `check_incl` scales logarithmically with the number of statements in the batch, we also ran it for varying numbers of statements. The results are in Figure 3.

7.2.4 Monitor overheads. Monitors must run a Bitcoin full node in order to get a complete uncensored view of the blockchain. As of December 2017, running a full node requires 145 GB of free disk space, increasing by up to 144 MB daily. It took us around three days to fully bootstrap a full node and verify all the blocks, although this operation needs to be performed only once per monitor.

7.2.5 Archival nodes overheads. Like monitors, archival nodes need to run a Bitcoin full node. Additionally, archival nodes must download and store all the data from the authority. The costs here are entirely dependent on the number and size of the statements; we examine the costs for Debian in Section 8.

In order for archival nodes to know which statement data to download from authorities to independently rebuild the Merkle tree roots committed in Bitcoin transactions and check that they match with the data provided, authorities must point the archival nodes to the location of the data. Again, this is dependent on the mechanism that the authority uses to make the data available.

As in Debian, however, archives use statements that represent files. We may therefore expect that, in addition to a Merkle tree, authorities would use metadata files to link each leaf in the tree to a file on the server that archival nodes then mirror; this would be particularly useful in a setting — like Debian — where it would be undesirable to reorganize files that are already stored. The metadata file would consist of a mapping of 32-byte hashes to filenames. The average Debian package filename is 60 bytes, so including such a metadata file would introduce an average storage overhead, for both authorities and archival nodes, of 92 bytes per statement.

8 USE CASE: DEBIAN

To demonstrate how Contour can be used on a real system, we prototyped it for auditing software binaries in the Debian software repository. Our results show that Contour provides a way to add transparency to this repository without major changes to the existing infrastructure and with minimal overheads. It could be deployed on top of the Debian ecosystem today, without any participant who did not want to opt in having to change their behavior.

We begin with an overview of how Debian currently works, and then go on to explain how existing actors in the ecosystem could play the roles necessary for Contour, along with the overheads.

8.1 Software distribution architecture

Debian is a popular Linux distribution used by over 32% of websites that run Linux.⁷ Software packages are installed and updated on Debian machines using the `apt` command-line program. The Debian software repository contains Release files for various versions of Debian, which are updated every time any package in the repository is updated. Each Release file contains a checksum for a Packages file, which contains a list of available software packages and their associated checksums for integrity checking.

Software packages are downloaded as `.deb` archives which provide the compiled binaries and scripts required to install a package on a system. These files are hosted in directories on HTTP mirrors, of which hundreds exist around the world.⁸

To cryptographically authenticate software packages, Debian has a set of tools called `apt-secure`. Debian installations come with a built-in set of PGP keys [15] that are used as trusted keys for validating software packages. Alongside the Release files in the repository, there are `Release.gpg` files that contain PGP signatures of the Release files under trusted PGP keys.⁹

Through the single signature of a Release file, `apt` can validate that individual `.deb` packages were authorised by a trusted PGP key by checking that the checksums of packages are included in the Packages file whose checksum is included in the root Release file. This of course creates a central point of failure, as the owner of the signing key can serve individual users targeted Release files — for example, if coerced to do so by law enforcement — that link to malicious packages.

8.2 Authority

In the case of Debian software distribution, the most natural operators for a Contour authority are the maintainers of the software repository. Specifically, the Contour authority would be the owner of the PGP key, as only this entity has the power to modify the software repository. Importantly, it is also possible for third parties to act as Contour authorities by proxy and commit binaries to the log on behalf of the maintainers of the Debian software repository. As committed binaries are transparent, the third party is not trusted any more than the maintainers of the Debian software repository would be, as any rogue additions to the log would still be detectable. This means it would be possible to deploy Contour today without any intervention or permission from the Debian project itself.

⁷<https://w3techs.com/technologies/details/os-linux/all/all>

⁸<https://www.debian.org/mirror/list>

⁹<https://wiki.debian.org/SecureApt>

To initiate the system as an authority, all the existing software packages would first need to be committed; i.e., the authority would need to commit to the current state of the repository. To measure the overhead needed for this step, we extracted the software package metadata for all processor architectures and releases of Debian from the Debian FTP archive (<https://www.debian.org/mirror/ftpmirror>) over a one-week period from January 20-27 2017. At the beginning of this period there were 976,214 unique software binaries available for download from the Debian software repositories, constituting 1.7 TB of data, and by the end there were 980,469.

As discussed above, the Debian package metadata already contains a SHA-256 hash for every packages, so we needed only to build a Merkle tree from these hashes (rather than compute them ourselves first), to then commit on the blockchain. This took approximately 6 seconds (which is in line with our benchmarks in Table 4 for 1M statements).

Going forward, the authority must commit batches of new and updated binaries to the log. The Debian FTP archives are updated four times a day, which means four batches to commit to the log per day. Recall from Section 7.2.2 that committing one transaction to the blockchain currently costs roughly 4.21 USD in fees, so this would cost 16.84 USD per day (although, as mentioned in Section 7, Bitcoin prices are notoriously volatile). This is a relatively low price to pay for a system that costs over 91M USD to attack (Section 6.1).

As the archive was updated, we kept track of the package hashes being added and created a new batch for each update. The average batch size was 1,040 packages, and the average time to build a Merkle tree for the batch was 0.0052 seconds.

As discussed in Section 7.2.5, we can also enable archival nodes to rebuild Merkle trees with minimal changes to the existing Debian archive infrastructure. This requires creating and storing only an additional 84 kB metadata file per batch, and an initial 79 MB metadata file. These metadata files consist of a mapping of hashes of software packages to their filenames in the Debian archives.

Finally, the proof of inclusion of each software package would need to be stored alongside each software package (.deb) file as metadata to be downloaded by Debian machines. At 980K software packages, this would require a maximum of 1.3 kB of extra storage per package, or 1.3 GB of extra storage to store the proofs of inclusion for all packages. Given the current storage requirements of (at least) 1.7 TB, this is only a 0.07% overhead.

8.3 Auditors

On the end-user side, the apt program would need to be modified to integrate the Auditor.check_incl and Auditor.sync calls, as implemented and analyzed in Section 7. This would ensure that downloaded packages are in the log before being installed.

In terms of overhead for end-user Debian machines, as discussed above this would require an extra 1.3 kB of bandwidth per package downloaded or updated. Given that the average package size is 1337 kB, the average overhead is 0.1% per package. We stress that this is a bandwidth requirement only, as once the proofs of inclusion are verified they do not need to be stored on the client's machine.

On a freshly installed Debian 8.8 system there are 520 packages installed by default, with a total .deb archive size of 190 MB. Verifying that each of these are in the log would require an extra 698.1 kB of bandwidth, and would take under two minutes.

8.4 Monitors

Debian's reproducible builds project allows any interested parties to verify that binaries published in the software repositories are compiled from a given source code.¹⁰ There are no specific parties assigned to the role of monitoring builds to see if they can be built from the source code. Similarly in Contour, any parties vested in the security of Debian may act as a monitor. Aside from end users, we anticipate that large organizations supplying critical infrastructure using Debian, national CERTs, and NGOs such as the Electronic Frontier Foundation would have an interest in monitoring the log.

Generally, any party that wants extra guarantees about the software updates they are installing — e.g., in order to be sure that the updates that have been pushed to their machines are the same as those that have been pushed to other machines — should run a monitor. For example, if a party running Debian receives update₁ and update₃ on their machine for some software package, but the log contains update₁, update₂, and update₃, then this raises a red flag as to why they did not receive update₂. In particular, update₂ may be a malicious update targeted to specific machines, and the party can check to see if the contents of update₂ have been made available by the authority. If they have not, then the authority is considered to be misbehaving. Optionally, honest archival nodes would prevent auditors from accepting the update altogether.

8.5 Archival nodes

There are 269 Debian mirrors hosting the full 1.7 TB archive, and we view these servers as the most natural candidates for operating archival nodes. The difference between a mirror and an archival node is that to fully satisfy availability an archival node should not delete any packages (even when packages are updated and removed), in order to enable monitors to examine obsolete packages.

In terms of overhead for archival nodes, this means storing the initial 1.7 TB, and then an additional average of 11 GB per day, or 4 TB per year. This is by far the highest overhead incurred by our system, and we expect that only a small number of mirrors would have the storage capacity to run an archival node. We stress that the use of archival nodes is optional and serves only to boost availability (as opposed to being required for integrity); moreover, there is currently at least one mirror hosting all historical Debian packages, so effectively already acting as an archival node¹¹.

8.6 Summary

In summary, Contour could be deployed on top of the existing system for Debian software distribution with minimal changes to the existing infrastructure. In terms of operating costs, the biggest overhead required to enable Contour is the extra storage space required for archival nodes (and again, this cost is optional). All other costs are minimal, with only a 0.07% storage overhead required for

¹⁰<https://wiki.debian.org/ReproducibleBuilds>

¹¹snapshot.debian.org/

the authority, and a 0.1% bandwidth overhead for the end user. The computational costs for these users are minimal as well.

One distinguishing feature of Contour is that no existing parties in the Debian infrastructure are required to participate if they do not want to, and as discussed earlier the security assumptions of the system would remain the same even if a third party acted as an authority. This places Contour in contrast to existing proposals for transparency (including some of the ones presented in Section 6.3), as they require the initial setup of some Sybil-free set of nodes. In contexts such as the distribution of Debian software packages, this assumption — and the security implications if it is violated — presents a significant obstacle to deployability, and avoiding this obstacle was one of our main goals in designing Contour.

9 DISCUSSION AND EXTENSIONS

Selective disclosure. When releasing software updates that patch critical security vulnerabilities, some software vendors may prefer not to reveal to potential attackers that, in the window of time in which a commitment has not yet been included in the blockchain, they can take advantage of victims with this vulnerable software installed. In such a case, Contour accounts for this by allowing the authority to commit to a batch of binaries visibly on the blockchain, but delay the publication of the binaries themselves until the commitment is sufficiently deep in the blockchain.

Generalized transparency. Although we have designed Contour for the specific application of binary transparency, the system is general enough to be applied to other applications requiring transparency. With the tradeoffs discussed in Section 6.3, it can even be applied to the setting of certificate transparency by using CAs as authorities, although it may be most beneficial in settings that present similar challenges to the ones discussed in the introduction (i.e., in which objects are large and persistent MitM attacks are a realistic threat).

Archival node scalability. The current design of Contour requires archival nodes to store all data, which as we have discussed in Section 8 incurs a significant overhead. There are likely many alternative designs that alleviate these requirements, such as a *sharded* solution in which archival nodes store only the data for which they sufficient space, and we leave an exploration of this space as an interesting open problem.

10 CONCLUSION

We have proposed Contour, a system that provides proactive transparency, logarithmic scaling for auditors in the number of packages they have installed, and does not require the initial coordination of forming a Sybil-free set of nodes. We have demonstrated that, even for attackers that are capable of performing persistent man-in-the-middle attacks, compromising the integrity of the system requires millions of dollars in energy and hardware costs. We also saw that Contour could be applied today to the Debian software repository with relatively low overhead to existing infrastructure, and with no changes or coordination required for any participant (even the Debian server) who does not wish to opt in.

REFERENCES

- [1] Security/binary transparency - mozillawiki, 2017. https://wiki.mozilla.org/Security/Binary_Transparency.
- [2] M. Apostolaki, A. Zohar, and L. Vanbever. Hijacking Bitcoin: Large-scale Network Attacks on Cryptocurrencies, 2016. arxiv.org/abs/1605.07524.
- [3] M. Bartoletti and L. Pompianu. An analysis of Bitcoin OP_RETURN metadata. In *4th Workshop on Bitcoin and Blockchain Research*, 2017.
- [4] D. Basin, C. Cremers, T. H.-J. Kim, A. Perrig, R. Sasse, and P. Szalachowski. ARPKI: Attack Resilient Public-Key Infrastructure. In *ACM CCS 2014*, pages 382–393, 2014.
- [5] J. Bonneau. EthKS: Using Ethereum to audit a CONIKS key transparency log. In *3rd Workshop on Bitcoin and Blockchain Research*, 2016.
- [6] M. Chase and S. Meiklejohn. Transparency Overlays and Applications. In *ACM SIGSAC Conference on Computer and Communications Security*, 2016.
- [7] L. Chuat, P. Szalachowski, A. Perrig, B. Laurie, and E. Messeri. Efficient Gossip Protocols for Verifying the Consistency of Certificate Logs. In *IEEE Conference on Communications and Network Security*, 2015.
- [8] X. de Carna  t de Carnavalet and M. Mannan. Challenges and implications of verifiable builds for security-critical open-source software. In *30th Annual Computer Security Applications Conference*, 2014.
- [9] B. Dowling, F. G  nther, U. Herath, and D. Stebila. Secure Logging Schemes and Certificate Transparency. In *ESORICS 2016*, 2016.
- [10] A. Eijdenberg, B. Laurie, and A. Cutter. Verifiable Data Structures, 2015. github.com/google/trillian/blob/master/docs/VerifiableDataStructures.pdf.
- [11] S. Eskandarian, E. Messeri, J. Bonneau, and D. Boneh. Certificate transparency with privacy. *CoRR*, abs/1703.02209, 2017.
- [12] I. Eyal and E. G. Sirer. Majority is not enough: Bitcoin mining is vulnerable. In *Financial Cryptography and Data Security*, 2014.
- [13] C. Farivar. Judge: Apple must help FBI unlock San Bernardino shooter’s iPhone, 2016. arstechnica.com/tech-policy/2016/02/judge-apple-must-help-fbi-unlock-san-bernardino-shooters-iphone/.
- [14] C. Fromknecht, D. Velicanu, and S. Yakubov. A decentralized public key infrastructure with identity retention. IACR Cryptology ePrint Archive, Report 2014/803, 2014. eprint.iacr.org/2014/803.pdf.
- [15] S. Garfinkel. *PGP: Pretty Good Privacy*. O’Reilly Media, Sebastopol, CA, USA, 1st edition, 1996.
- [16] A. Gervais, H. Ritzdorf, G. Karame, and S. Capkun. Tampering with the Delivery of Blocks and Transactions in Bitcoin. In *ACM CCS 2015*, 2015.
- [17] D. Goodin. “Flame” malware was signed by rogue Microsoft certificate, 2012. arstechnica.com/security/2012/06/flame-malware-was-signed-by-rogue-microsoft-certificate/.
- [18] E. Heilman, A. Kendler, A. Zohar, and S. Goldberg. Eclipse Attacks on Bitcoin’s Peer-to-Peer Network. In *USENIX Security 2015*, 2015.
- [19] T. H.-J. Kim, L.-S. Huang, A. Perrig, C. Jackson, and V. Gligor. Accountable key infrastructure (AKI): a proposal for a public-key validation infrastructure. In *WWW 2013*, pages 679–690, 2013.
- [20] E. K. Kogias, P. Jovanovic, N. Gailly, I. Khoffi, L. Gasser, and B. Ford. Enhancing Bitcoin Security and Performance with Strong Consistency via Collective Signing. In *USENIX Security 2016*, 2016.
- [21] B. Laurie, A. Langley, and E. Kasper. Certificate transparency, 2013.
- [22] J. Leyden. Inside ‘Operation Black Tulip’: DigiNotar hack analysed, 2011. www.theregister.co.uk/2011/09/06/diginotar_audit_damning_fail/.
- [23] S. Matsumoto and R. M. Reischuk. IKP: Turning a PKI Around with Blockchains. IACR Cryptology ePrint Archive, Report 2016/1018, 2016. eprint.iacr.org/2016/1018.
- [24] M. S. Melara, A. Blankstein, J. Bonneau, E. W. Felten, and M. J. Freedman. CONIKS: Bringing Key Transparency to End Users. In *USENIX Security 2015*, 2015.
- [25] S. Nakamoto. Bitcoin: A Peer-to-Peer Electronic Cash System, 2008. bitcoin.org/bitcoin.pdf.
- [26] K. Nikitin, E. Kokoris-Kogias, P. Jovanovic, N. Gailly, L. Gasser, I. Khoffi, J. Cappsos, and B. Ford. CHAINIAC: Proactive software-update transparency via collectively signed skipchains and verified builds. In *26th USENIX Security Symposium (USENIX Security 17)*, pages 1271–1287, Vancouver, BC, 2017. USENIX Association.
- [27] L. Nordberg, D. Gillmor, and T. Ritter. Gossiping in CT, 2016. tools.ietf.org/html/draft-ietf-trans-gossip-03.
- [28] M. D. Ryan. Enhanced Certificate Transparency and End-to-end Encrypted Mail. In *NDSS 2014*, 2014.
- [29] A. Singh, T.-W. J. Ngan, P. Druschel, and D. S. Wallach. Eclipse attacks on overlay networks: Threats and defenses. In *IEEE Conference on Computer Communications*, 2006.
- [30] E. Syta, I. Tamas, D. Visher, D. I. Wolinsky, P. Jovanovic, L. Gasser, N. Gailly, I. Khoffi, and B. Ford. Keeping Authorities “Honest or Bust” with Decentralized Witness Cosigning. In *IEEE Symposium on Security and Privacy (“Oakland”)*, 2016.
- [31] A. Tomescu and S. Devadas. Catena: Efficient Non-equivocation via Bitcoin. In *IEEE Symposium on Security and Privacy (“Oakland”)*, 2017.




Figure 4: The number of Antminer S9 rigs required to produce blocks under a certain time limit.

A COST OF A SPLIT-VIEW ATTACK

To support our argument in Section 6.1 about the infeasibility of carrying out a split-view attack, we provide here more concrete estimates for the associated costs of the attack. These are rough estimates, as they make assumptions about certain properties (e.g., electricity costs and choice of mining hardware) that are not guaranteed to hold in practice. We are not aware of any previous literature considering the costs of eclipse attacks on Bitcoin nodes, so we consider these estimates (even if rough) to be important.

We first calculate the cost to mine a single block, and then analyze the cost of performing a split-view attack in the case where the adversary is able to perform an eclipse attack and where it cannot.

Cost to mine a single block. The probability of a miner finding a valid block after each hashing attempt is $\frac{2^{16}-1}{2^{48}D}$, where D is the periodically adjusted difficulty of the network. For a miner to mine a block then, they must make on average $\frac{2^{48}D}{2^{16}-1}$ hashing attempts. The total electricity cost (C) of mining a block is thus

$$C = \frac{2^{48}D}{2^{16}-1} \cdot J \cdot E, \quad (1)$$

where J is the number of joules required per hashing attempt, and E is the electricity cost of one joule. As of December 2017, the most energy-efficient Bitcoin mining hardware is the Antminer S9, which has an energy cost of $9.82 \cdot 10^{-11}$ joules per hash,¹² and the average retail price of one kilowatt hour in the US is 0.10 USD.¹³ The cost per joule, E , is therefore $\frac{0.10}{1000 \cdot 60 \cdot 60} = 2.8 \cdot 10^{-8}$ USD. As of December 2017, the Bitcoin mining difficulty (D) is 1,347,001,430,558. Plugging these numbers into Equation 1, the total electricity cost to mine a block, using the most efficient hardware and assuming standard electricity costs, is thus 15,908 USD.

¹²en.bitcoin.it/wiki/Mining_hardware_comparison

¹³www.eia.gov/electricity/state/

To also take hardware costs into account, the number of mining rigs N needed to mine a block in S seconds is

$$N = \frac{\left(\frac{2^{48}D}{2^{16}-1}\right)}{H \cdot S}, \quad (2)$$

where H is the number of hashes that the mining rig is capable of calculating per second. This formula is graphed in Figure 4 for the Antminer S9 rig, which is capable of calculating 14 terahashes per second and has a retail cost of 2,400 USD.¹⁴ We use these formulas to estimate the cost of split-view attacks in the following analysis.

Using eclipse attacks. If an eclipse attack is possible, an adversary can launch a successful split-view attack solely by mining k blocks at its own pace, where k is the number of blocks the auditor requires to be mined after a block containing a given commitment in order to consider that commitment as valid. (It is standard in most Bitcoin wallets to use $k = 6$.)

Using our rough estimates above, it would cost the adversary 15,908 USD in electricity costs to mine a block, or 95,448 USD for $k = 6$. The hardware costs depends on how much time the adversary needs to conduct the attack, or how long they are able to continue their man-in-the-middle attack on the auditor. If — as a conservative number — the adversary wants to conduct the attack within a week, it must mine a block every 1.4 days to produce 6 blocks, which requires 3,417 mining rigs at a hardware cost of 8,200,800 USD. This brings the total cost of the attack to 8.3M USD. Moreover, this attack is also fundamentally targeted: if the adversary wants to later compromise previously non-eclipsed auditors, it must mine a new set of blocks (assuming these auditors have more up-to-date blocks) and pay the electricity costs again. Even for an adversary with few financial constraints, this makes it significantly more difficult to conduct such an attack on a wide scale.

Furthermore, if the adversary takes 1.4 days to mine a block, or in general the auditor sees no new blocks until long after the expected 10-minute interval, it may assume that an eclipse attack is being performed. We can thus greatly increase the cost of the attack by adding simple checks to the auditor to ensure that there is a maximum interval between blocks. If we generously set such a check to require a maximum of 3 hours between blocks, then a total of 38,263 mining rigs are required at a cost of 91.8M USD.

In addition, the blocks must still follow the same difficulty level as honest blocks, so by mining these only in the eclipsed view of the network the adversary is not only expending the energy needed to do so but is also forfeiting the mining reward associated with them. As of December 5 2017, the Bitcoin mining reward is 12.5 bitcoins, or roughly 145,250 USD, so for $k = 6$ the adversary must additionally forfeit 871,500 USD.

Ignoring eclipse attacks. To perform a split-view attack without an eclipse attack, an adversary must fork the Bitcoin blockchain, which naïvely requires control of 51% of the network’s mining power.

As of December 5 2017, the total hashing power of the Bitcoin network was 11,918,845 terahashes per second.¹⁵ Conducting a 51% attack would therefore require the adversary to be able to compute

¹⁴www.amazon.com/Antminer-S9-0-10W-Bitcoin-Miner/dp/B01GFEOV00

¹⁵blockchain.info/charts/hash-rate

more than 11,918,845 terahashes per second. Per hour, the total electricity cost would be $11918845 \cdot 10^{12} \cdot 3600 \cdot J \cdot E$, or — using our earlier estimates for J and E — 117,979 USD per hour. In terms of hardware costs, if we use the figures for the Antminer S9 from before, the total number of mining rigs required would be greater than $\frac{11918845 \cdot 10^{12}}{14 \cdot 10^{12}} = 851346$, at a total cost of 2043M USD.

While more sophisticated attacks, such as selfish mining [12], have proposed strategies that fork the blockchain using only 25% of the mining power, this would still require an investment of hundreds of millions of dollars. Such an attack would furthermore be highly visible, as the blockchain is regularly monitored for forks.