

# Contour: A Decentralized System for Binary Transparency

## Abstract

Transparency is crucial in security-critical applications that rely on authoritative information, as it provides a robust mechanism for holding these authorities accountable for their actions. A number of solutions have emerged in recent years that provide transparency in the setting of certificate issuance, and Bitcoin provides an example of how to enforce transparency in a financial setting. In this work, we shift to a new setting, the distribution of software package binaries, and present a system for so-called “binary transparency.” Our solution, Contour, uses proactive methods for providing transparency, privacy, and availability, even in the face of persistent man-in-the-middle attacks. We also demonstrate, via benchmarks and a test deployment for the Debian software repository, that Contour is efficient enough to be used in practice and requires minimal coordination effort to deploy.

## 1 Introduction

Historically, functional societies have relied to a large degree on trust in their governing institutions, with participants in various systems (nation states, the Internet, financial markets, etc.) trusting those in charge to follow an agreed-upon set of rules and thus provide the system with some level of integrity. In recent years, however, increasing numbers of incidents have demonstrated that integrity cannot be meaningfully achieved solely by placing trust in a small number of entities. As a result, people are now demanding more active participation in the systems with which they interact, and more accountability for the entities that govern them. One of the main methods that has been relatively successful in achieving accountability is the idea of *transparency*, in which information about the decisions within the system are made globally visible, thus enabling any participant to check for themselves whether or not the decisions seem to be

complying with what they perceive to be the rules.

One of the technical settings in which the idea of transparency has been most thoroughly — and successfully — deployed is the issuance of X.509 certificates; this is likely due partially to the nature of these certificates (which are themselves intended to be globally visible), and partially to the many publicized failures of major certificate authorities (CAs) [14, 19]. A long line of recent research [18, 3, 16, 25, 21, 8, 20, 28] has provided and analyzed solutions that bring transparency to the issuance of both X.509 certificates (“certificate transparency”) and to the assignment of public keys to end users (“key transparency”).

Despite their differences, many of these systems share a fundamentally similar architecture [6]: after being signed by CAs, certificates are stored by *log servers* in a globally visible append-only log; i.e., in a log in which entries cannot be deleted without detection. Clients are told to not accept certificates unless they have been included in such a log, and to determine this they rely on *auditors*, who are responsible for checking inclusion of the specific certificates seen by clients. Because auditors are often thought of software running on the client (e.g., a browser extension), they must be able to operate efficiently. Finally, in order to expose misbehavior, *monitors* (inefficiently) inspect the certificates stored in a given log to see if they satisfy the rules of the system.

To prevent clients from accepting bad certificates, such systems thus rely on the monitors to expose them. Because auditors are the ones communicating with the client, however, to achieve this property an additional line of communication is needed between the auditor and monitor in the form of a *gossip* protocol [24, 7]. In such a protocol, the auditor and monitor exchange information on their current view of the log, which allows them to detect whether or not their views are *consistent*, and thus whether or not the log server is misbehaving by presenting “split” views of the log. If such attacks are possible, then the accountability of the system is destroyed, as a

log server can present one log containing all certificates to auditors (thus convincing it that its certificates are in the log), and one log containing only “good” certificates to monitors (thus convincing them that all participants in the system are obeying the rules).

While gossiping can detect this misbehavior, it is ultimately a retroactive mechanism — i.e., it detects rather than prevents this misbehavior — and is thus most effective in settings where (1) no man-in-the-middle (MitM) attack can occur, so the line of communication between an auditor and monitors remains open, and (2) some form of external punishment is possible, to sufficiently disincentivize misbehavior on the basis of detection. Various systems have been proposed recently that are designed to operate in settings where these assumptions cannot be made, such as Collective Signing [27], but perhaps the most prominent example of such a system is Bitcoin (and all cryptocurrencies based on the idea of a *blockchain*). In Bitcoin, all participants have historically played the simultaneous role of log servers (in storing all Bitcoin transactions), auditors, and monitors (in checking that no double-spending takes place). The high level of integrity achieved by this comes at great expense to the participants, both in terms of storage costs (the Bitcoin blockchain is currently over 100 GB<sup>1</sup>) and computational costs in the form of the expensive proof-of-work mechanism required to maintain the blockchain, but several recent proposals attempt to achieve the same level of integrity in a more scalable way [28, 17].

Because of the effectiveness of these approaches, there has been interest in repurposing them to provide not only transparency for certificates or monetary transfers, but for more general classes of objects (“general transparency” [9]). One specific area that thus far has been relatively unexplored is the setting of software distribution (“binary transparency”<sup>2</sup>). Bringing transparency to this setting is increasingly important, as there are an increasing number of cases in which actors target devices with malicious software signed by the authoritative keys of update servers. For example, the Flame malware, discovered in 2012, was signed by a rogue Microsoft certificate [14] and masqueraded as a routine Microsoft software update [23]. In 2016, a US court compelled Apple to produce and sign custom firmware in order to disable security measures on a phone that the FBI wanted to unlock [11].

Aside from its growing relevance, binary transparency is particularly in need of exploration because the techniques described above for both certificate transparency and Bitcoin cannot be directly translated to this setting. Whereas certificates and Bitcoin transactions are small

(on the order of kilobytes), software binaries can be arbitrarily large (often on the order of gigabytes), so cannot be easily stored and replicated in a log or ledger. More importantly, by their very nature software packages have the ability to execute code on a system, so malicious software packages can easily disable gossiping mechanisms or launch wider (and long-lasting) MitM attacks. This makes retroactive methods for detecting misbehavior almost uniquely poorly suited to this setting, in which clients need to know that a software package has been inspected by independent parties *before* installing it, not after.

**Our contributions.** We present Contour, a solution for binary transparency that utilizes the Bitcoin blockchain in an efficient manner to proactively prevent clients from installing malicious software, even in the face of long-term MitM attacks.

We begin in Section 4 by presenting our threat model. In addition to the goal of preventing split views, we highlight the importance of *auditor privacy*, in which auditors should not reveal the particular binaries in which they are interested (as this could reveal, for example, that a client has a susceptible version of some software), and of *availability*, in which auditors and monitors should still be able to do their job even if the original software update server loses its data or goes offline.

After then presenting the design of Contour in Section 5, we go on to analyze both its security and its efficiency in Section 6. Given the volume of related research on certificate transparency, we also present some comparisons here, and argue that ours is the first efficient solution to provide these security guarantees without requiring any coordination cost, in the form of selecting a central entity to perform authorization, or otherwise forming a Sybil-free set of log servers. This ensures that Contour could be easily and securely deployed today.

To validate our efficiency claims, in Section 7 we describe an implementation of Contour and benchmark its performance, finding that almost all operations can be performed very quickly (on the order of microseconds), that auditors can store minimal information (on the order of kilobytes), and that arbitrary numbers of binaries can be represented by a single small (235-byte) Bitcoin transaction. We also validate our claims of real-world relevance by presenting, in Section 8, the application of Contour to the current package repository for the Debian operating system. We find that it would require minimal overhead for the existing actors within this system, and cost under 2 USD per day.

Finally, in Section 9 we present some possible extensions to Contour, including a discussion of how to use it to achieve general transparency, and in Section 10 we conclude.

<sup>1</sup>[blockchain.info/charts/blocks-size](http://blockchain.info/charts/blocks-size)

<sup>2</sup>[groups.google.com/forum/#!forum/binary-transparency](https://groups.google.com/forum/#!forum/binary-transparency)

## 2 Related Work

As mentioned in the introduction, there is at this point a significant volume of related work on the idea of transparency, particularly in the settings of certificates, keys, and Bitcoin. We briefly describe some of this work here, and provide a more thorough comparison to the most relevant work in Section 6.3. To the best of our knowledge, ours is the first full working solution in the context of binary transparency.

In terms of certificate and key transparency, AKI [16] and ARPKI [3] provide a distributed infrastructure for the issuance of certificates, thus providing a way to prevent rather than just detect misbehavior. Certificate Transparency (CT) [18] focuses on the storage of certificates rather than their issuance, Ryan [25] demonstrated how to handle revocation within CT, and Dowling et al. [8] provided a proof of security for it. CONIKS [21] focuses instead on key transparency, and thus pays more attention to privacy and does not require the use of monitors (but rather has users monitor their own public keys).

In terms of solutions that avoid gossip by using Bitcoin (or more generally using a blockchain), Fromknecht et al. [12] propose a decentralized PKI based on Bitcoin and Namecoin, and IKP [20] provides a method for certificate issuance based on Ethereum. EthIKS [4] provides an Ethereum-based solution for key transparency and, concurrently with our work, Catena [28] provides one based on Bitcoin. While both Catena and Contour utilize similar recent features of Bitcoin to achieve efficiency, they differ in their focus (key vs. binary transparency), and thus in the proposed threat model; e.g., they dismiss eclipse attacks [26] on the Bitcoin network, whereas we consider them well within the scope of a MitM attacker.

Finally, in terms of more general solutions, Chase and Meiklejohn abstract CT into the general idea of a “transparency overlay” [6] and prove its security. Similarly, Collective Signing [27] is a general consensus mechanism that shares our goal of providing transparency even in the face of MitM attacks and thus avoids gossiping, but requires setting up a distributed set of “witnesses” that is free of Sybils. This is a deployment overhead that we seek to avoid.

## 3 Background

### Distributed ledgers

The concept of a blockchain was first used in Bitcoin, which is designed to be a globally consistent append-only ledger of financial transactions, and was introduced by (the pseudonymous) Satoshi Nakamoto [22]. Given our limited usage of Bitcoin, we ignore the thousands of other cryptocurrencies that have now been introduced,

and for brevity we focus only on the properties of Bitcoin that we require for Contour. For further discussion of the more general properties of the blockchain, we point the interested reader to Bonneau et al. [5].

Briefly, the Bitcoin blockchain is literally a chain of blocks. Each block contains two components: a *header* and a list of transactions. In addition to other metadata, the header stores the hash of the block (which, in compliance with the proof-of-work consensus mechanism, must be below some threshold in order to show that a certain amount of so-called “hashing power” has been expended to form the block), the hash of the previous block (thus enabling the chain property), and the root of the Merkle tree that consists of all transactions in the block.

On the constructive side, while the scripting language used by Bitcoin is (intentionally) limited in its functionality, Bitcoin transactions can nevertheless store small amounts of arbitrary data. This makes Bitcoin potentially useful for other applications that may require the properties of its ledger, such as certifying the ownership and timestamp of a document [2]. One mechanism that allows Bitcoin to store such data is the script opcode `OP_RETURN`,<sup>3</sup> which can be used to embed up to 80 bytes of arbitrary data.

Another aspect of Bitcoin that enables additional development is the idea of an SPV (Simplified Payment Verification) client.<sup>4</sup> Rather than perform the expensive verification of the digital signatures contained in Bitcoin transactions, or the checks necessary to determine whether or not double-spending has taken place, these clients check only that a given transaction has made it into some block in the blockchain. As this can be achieved using only the root hashes stored in the block headers, such clients can store only these headers (which are small) and verify only Merkle proofs of inclusion (which is fast), and are thus significantly more efficient than their “full node” counterparts.

On the destructive side, various attacks have been demonstrated that undermine the security guarantees of Bitcoin. In *eclipse* attacks [15, 13, 1], an adversary exploits the topology of the Bitcoin network to interrupt, or at least delay, the delivery of announcements of new transactions and blocks to a victim node. More expensive “51%” attacks, in which the adversary controls more than half of the collective hashing power of the network, allow the adversary to fork the blockchain, and it has been demonstrated [10] that such attacks can in fact be carried out with far less than 51% of the hashing power.

<sup>3</sup>[en.bitcoin.it/wiki/OP\\_RETURN](https://en.bitcoin.it/wiki/OP_RETURN)

<sup>4</sup>[en.bitcoin.it/wiki/Scalability#Simplified\\_payment\\_verification](https://en.bitcoin.it/wiki/Scalability#Simplified_payment_verification)

## Software distribution

Software distribution on modern desktop and mobile operating systems is managed through centralized software repositories such as the Apple App Store, the Android Play Store, or the Microsoft Store. Most Linux distributions such as Debian also have their own software repositories from which administrators can install and update software packages using the `apt` command-line program.

To tackle the issue of transparency in these repositories, *reproducible builds* allow users to verify that a compiled binary matches a version of the source code. This is, however, a resource-expensive manual verification that standard users are not expected to do, and does not guarantee that the source code is the same source code that everyone else has received.

## 4 Threat Model and Setting

In this section, we describe the actors in our transparency overlay for the software distribution ecosystem (Section 4.1), along with the interactions between these actors (Section 4.2), and the goals we hope to achieve by using such an overlay (Section 4.3).

### Participants

We consider a system with five types of actors: services, authorities, monitors, auditors, and clients. We describe each of these types below in the singular, but for the correct and secure functioning of a transparency overlay we require a distributed set of auditors and monitors, each acting independently.

**Service:** The service is responsible for producing actions, such as the issuance of a software update. In order to have these actions authorized, they must be sent to the *authority*.

**Authority:** The authority is responsible for publishing *statements* that declare it has seen a given action taken by a service; e.g., that it has been sent a given software binary. These statements furthermore claim that the authority has—in some form—published these actions in a way that allows them to be inspected by the *monitor*. As this inspection is typically inefficient, the authority is also responsible for placing its statements into a public *audit log*, where they can be efficiently verified by the *auditor*.

**Monitor:** The monitor is responsible for inspecting the actions published by the authority and performing out-of-band tests to determine their validity (e.g., to ensure that software updates do not contain malware).

**Auditor:** The auditor is responsible for checking specific actions against the statements made by the authority that claim they are published.

**Client:** The client receives software updates from either the authority or the service, along with a statement that claims the update has been published for inspection. It outsources all responsibility to the auditor, so in practice the auditor can be thought of as software that sits on the client (thus making the client and auditor the same actor).

### Interactions

In terms of the interactions between these entities, one of the main benefits of Contour—as discussed in the introduction—is that entities do not need to engage in prolonged multi-round interactions like gossiping, but rather pass messages atomically to one another. As we see in Section 6.1, this minimizes the advantage that an adversary could gain by launching man-in-the-middle attacks. We therefore outline only non-interactive algorithms needed to generate messages, rather than interactive protocols, and wait to specify the exact inputs and outputs until we present our construction in Section 5.

**Authority.commit:** The authority runs this algorithm to commit statements to the audit log.

**Authority.prove\_incl:** The authority runs this algorithm to provide a proof that a specific statement is in the audit log.

**Auditor.check\_incl:** The auditor runs this algorithm to check the proof of inclusion for a specific statement.

**Monitor.get\_commits:** The monitor runs this algorithm to retrieve relevant commitments from the audit log.

### Goals

We break the goals of the system down into security goals (denoted with an S) and deployability goals (denoted with a D).

In all our security goals, we aim to defend against the specified attacks in the face of malicious authorities that, in addition to performing all the usual actions of the authority, can also perform man-in-the-middle attacks on the auditor’s network communications, and can compromise the client’s machine. If additional adversaries are considered we state them explicitly.

**S1: No split views.** We would like to prevent split-view attacks, in which the information contained in the audit log convinces the auditor that the authority published an action taken in the system, and thus it is able

to be inspected by monitors, whereas in fact it is not and only appears that way in the auditor’s “split” view of the log.

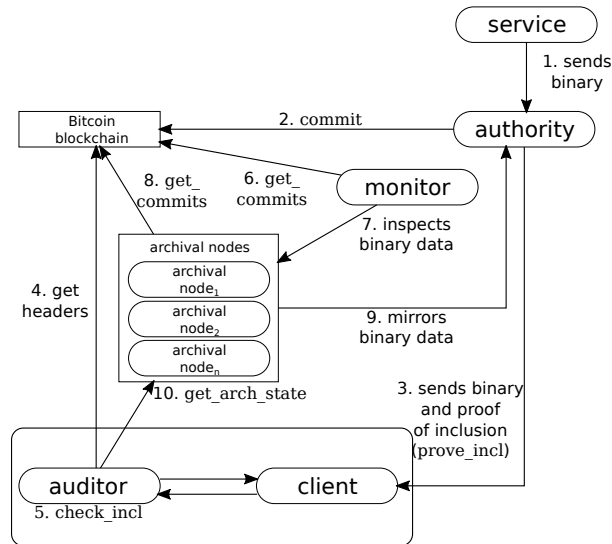
**S2: Availability.** We would like to prevent attacks on availability, in which the information contained in the audit log convinces the auditor that an action is available to be inspected by monitors, when in fact the authority has not published it or has, after the initial publication, lost it or intentionally taken it down.

**S3: Auditor privacy.** We would like to ensure that the specific binaries in which the auditor is interested are not revealed to any other parties (as this might reveal, for example, that a client has a software version that is susceptible to malware). We thus consider how to achieve this not only in the face of malicious authorities, but in the case in which all parties aside from the auditor are malicious.

**D1: Efficiency.** We would like Contour to operate as efficiently as possible, in terms of computational, storage, and communication costs. In particular, we would like the overhead beyond the existing requirements for a software distribution system to be minimal.

**D2: Minimal setup.** In addition to the computational overheads, we would like as little effort — in terms of, e.g., coordination — to be done as possible in order to deploy Contour.

## 5 Design of Contour



**Figure 1:** The overall structure of Contour.

In this section we describe the overall design of Con-

tour. An overview of the interactions between all the various actors can be seen in Figure 1.

### Setup and instantiation

Contour and its security properties makes use of a blockchain, whose primary purpose — as we see in Section 6.1 — is to provide an immutable ledger that prevents split-view attacks. Because the Bitcoin blockchain is currently the most expensive to attack, we use it here and in our security analysis in Section 6.1, but observe that any blockchain could be used in its place. As a result of using Bitcoin, every authority must initially establish a known Bitcoin address, which we can think of as a hard-coded value in the auditor software.

### Logging and publishing statements

To start, the authority receives actions from services; i.e., software binaries from the developers of the relevant packages (Step 1 of Figure 1). As it receives such a binary, it incorporates its hash as a leaf in a Merkle tree with root  $h_T$ . The root, coupled with the path down to the leaf representing the binary, thus proves that the authority has seen the binary, so we view the root as a batched statement attesting to the fact that the authority has seen all the binaries represented in the tree. Once the Merkle tree reaches some (dynamically chosen) threshold  $n$  in size, the authority runs the `commit` algorithm (Step 2 of Figure 1) as follows:

`commit( $h_T$ )`: Form a Bitcoin transaction in which one of the outputs embeds  $h_T$  by using `OP_RETURN`, and the other inputs and outputs are chosen according to the authority’s available addresses and the fees it wants to pay. (See Section 7.2 for some concrete choices.) Publish the transaction to the Bitcoin blockchain and return the raw transaction data, denoted  $tx$ .

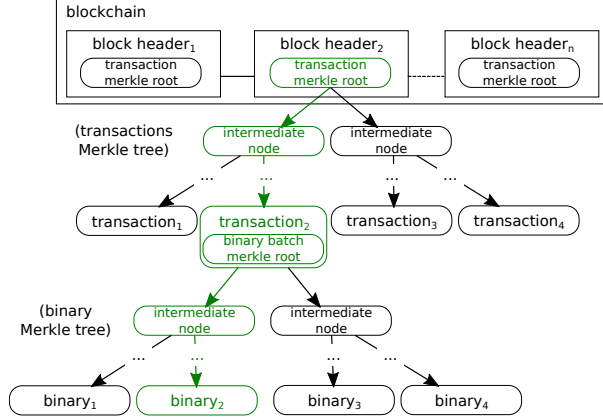
Crucially, the `commit` algorithm stores only the root hash in the transaction, meaning its size is independent of the number of statements it represents. Furthermore, if the blockchain is append-only — i.e., if double spending is prevented — then the log represented by the commitments in the blockchain is append-only as well.

### Proving inclusion

After committing a batch of binaries to the blockchain, the authority can now make these binaries accessible to clients. When a client requests a software update, the authority sends not only the relevant binary, but also an accompanying proof of inclusion, which asserts that the binary has been placed in the log and is thus accessible to monitors (Step 3 of Figure 1).

To generate this proof, the authority must first wait for its transaction to be included in the blockchain (or, for

improved security, for it to be embedded  $k$  blocks into the chain). We denote the header of the block in which it was included as  $\text{head}_B$ . The proof then needs to convince anyone checking it of two things: (1) that the relevant binary is included in a Merkle tree produced by the authority and (2) that the transaction representing this Merkle tree is in the blockchain. Thus, as illustrated in Figure 2, this means providing a path of hashes leading from the values retrieved from the blockchain to a hash of the statement itself.



**Figure 2:** An example of a path of hashes leading from the block’s transactions merkle root to the hash of the statement<sub>2</sub>.

For a given binary  $\text{binary}$ , the algorithm `prove_incl` thus runs as follows:

`prove_incl(tx, headB, binary)`: First, form a Merkle proof for the inclusion of  $\text{tx}$  in the block represented by  $\text{head}_B$ . This means forming a path from the root hash stored in  $\text{head}_B$  to the leaf representing  $\text{tx}$ ; denote these intermediate hashes by  $\pi_{\text{tx}}$ . Second, form a Merkle proof for the inclusion of  $\text{binary}$  in the Merkle tree represented by  $\text{tx}$  (using the hash  $h_T$  stored in the `OP_RETURN` output) by forming a path from  $h_T$  to the leaf representing  $\text{binary}$ ; denote these intermediate hashes by  $\pi_{\text{binary}}$ . Return  $(\text{head}_B, \text{tx}, \pi_{\text{tx}}, \pi_{\text{binary}})$ .

## Verifying inclusion

To verify this proof, the auditor must check the Merkle proofs, and must also check the authority’s version of the block header against its own knowledge of the Bitcoin blockchain. This means that the auditor must first keep up-to-date on the headers in the blockchain, which it obtains by running an SPV client (Step 4 in Figure 1). By running this client, the auditor builds up a set  $S = \{\text{head}_{B_i}\}_i$  of block headers, which it can check against the values in the proof of inclusion. This means that, for a binary  $\text{binary}$ , `check_incl` (Step 5 in Figure 1) runs as follows:

`check_incl(S, binary, (headB, tx,  $\pi_{\text{tx}}$ ,  $\pi_{\text{binary}}$ ))`: First, check that  $\text{head}_B \in S$ ; output 0 if not. Next, extract  $h_T$  from  $\text{tx}$  (using the hash stored in the `OP_RETURN` output), form  $h_{\text{binary}} \leftarrow H(\text{binary})$ , and check that  $\pi_{\text{binary}}$  forms a path from the leaf  $h_{\text{binary}}$  to the root  $h_T$ . Finally, form  $h_{\text{tx}} \leftarrow H(\text{tx})$ , and check that  $\pi_{\text{tx}}$  forms a path from the leaf  $h_{\text{tx}}$  to the root hash in  $\text{head}_B$ . If both these checks pass then output 1; otherwise output 0.

## Ensuring availability

Independently of auditors, monitors must retrieve all commitments associated with the authority from the blockchain and mirror their binaries (Steps 6 and 7 of Figure 1). This means `get_commits` runs as follows:

`get_commits()`: Retrieve all transactions in the blockchain associated with the authority’s address(es), and return the hashes stored in the `OP_RETURN` outputs.

After checking the binaries against their commitments, the monitors then inspect them — to, e.g., ensure they are not malware — in ways we consider outside of the scope of this paper.

While the system we have described thus far functions correctly, in order to make binaries available for inspection, we assume the monitors can mirror the authority’s logs. It therefore fails to satisfy our goal of availability in the (realistic) event that the authority goes down at some point in time.

We thus consider the case where the authority commits binaries to the blockchain, but — either intentionally or because it loses the data sometime in the future — does not supply the data to monitors. While this is detectable by monitors, as they can see that there are commitments in the blockchain with no data behind them, to disincentive this behavior requires some retroactive real-world method of punishment. More importantly, it prevents the monitor from pinpointing specific bad actions, such as malicious binaries, and thus from identifying the potential victims of the authority’s misbehavior.

Because of this, it is thus desirable to not only enable the detection of this form of misbehavior, but in fact to prevent it from happening in the first place. One way to achieve this is to have auditors mirror the binary themselves and send it to monitors before accepting it, to ensure that they have seen it and believe it to be benign. While this would be effective, and is arguably practical in a setting such as Certificate Transparency (modulo concerns about privacy) where the objects being sent are relatively small, in the setting of software distribution — where the objects being sent are binaries, and are thus quite large — it is too inefficient to be considered.

Instead, we propose a new actor in the ecosystem presented in Section 4: archival nodes, or *archivists*, that are responsible for mirroring all data from the authority (Steps 8 and 9 in Figure 1). To gain the extra guarantee that the data is available to monitors, auditors may thus use any archival nodes of which they are aware to check their state (i.e., the most recent block header for which they have data from the authority) and ensure that they cover the block headers relevant to the proofs they are checking (Step 10 in Figure 1). This means adding the following two algorithms to the list in Section 4.2:

`Archivist.get_commits()`: The archivist runs this algorithm to access the commitments made by the authority, just as is done by the monitor (using the same algorithm).

`Auditor.get_arch_state()`: The auditor (optionally) runs this algorithm to obtain the state of any archivists of which it is aware. This is simply the latest block header for which the archival node has mirrored the data behind the commitments held within.

## 6 Evaluation

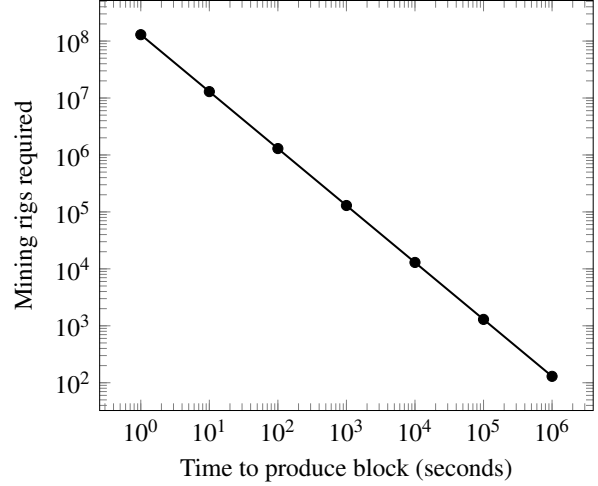
In this section, we evaluate Contour in terms of how well it meets the security goals (Section 6.1) and deployability goals (Section 6.2) specified in our threat model in Section 4.3. We also compare it with respect to previous solutions in Section 6.3, and argue that it is the only system to achieve all our goals.

### Security goals

#### No split views (S1)

In order to prevent split views, we rely on the security of the Bitcoin blockchain and its associated proof-of-work-based consensus mechanism. If every party has the same view of the blockchain, then split views of the log are impossible, as there is a unique commitment to the state of the log at any given point in time. The ability to prevent split views therefore reduces to the ability to carry out attacks on the Bitcoin blockchain. We break down the costs of such attacks below: in particular, we first consider the cost of mining a single block, and then separately examine the case when the adversary can carry out an eclipse attack — in which, recall from Section 3, it can control the auditor’s view of the blockchain — and the case when it cannot.

**Cost to mine a single block.** The probability of a miner finding a valid block after each hashing attempt is  $\frac{2^{16}-1}{2^{48}D}$ , where  $D$  is the periodically adjusted difficulty of the network. For a miner to mine a block then, they must make on average  $\frac{2^{48}D}{2^{16}-1}$  hashing attempts. The total



**Figure 3:** The number of Antminer S9 rigs required to produce blocks under a certain time limit.

electricity cost ( $C$ ) of mining a block is thus

$$C = \frac{2^{48}D}{2^{16}-1} \cdot J \cdot E, \quad (1)$$

where  $J$  is the number of joules required per hashing attempt, and  $E$  is the electricity cost of one joule. As of February 2017, the most energy-efficient Bitcoin mining hardware is the Antminer S9, which has an energy cost of  $9.82 \cdot 10^{-11}$  joules per hash,<sup>5</sup> and the average retail price of one kilowatt hour in the US is 0.10 USD.<sup>6</sup> The cost per joule,  $E$ , is therefore  $\frac{0.10}{1000 \cdot 60 \cdot 60} = 2.8 \cdot 10^{-8}$  USD. As of February 2017, the Bitcoin mining difficulty ( $D$ ) is 422,170,566,884. Plugging these numbers into Equation 1, the total electricity cost to mine a block, using the most efficient hardware and assuming standard electricity costs, is thus 4,986 USD.

To also take account of hardware costs, we observe that the number of mining rigs  $N$  needed to mine a block in  $S$  seconds is

$$N = \frac{\left(\frac{2^{48}D}{2^{16}-1}\right)}{H \cdot S}, \quad (2)$$

where  $H$  is the number of hashes that the mining rig is capable of calculating per second. This formula is graphed in Figure 3 for the Antminer S9 rig, which is capable of calculating 14 terahashes per second and has a retail cost of 2,400 USD.<sup>7</sup> We use these formulas to estimate the cost of split-view attacks in the following analysis.

<sup>5</sup>[en.bitcoin.it/wiki/Mining\\_hardware\\_comparison](http://en.bitcoin.it/wiki/Mining_hardware_comparison)

<sup>6</sup>[www.eia.gov/electricity/state/](http://www.eia.gov/electricity/state/)

<sup>7</sup>[www.amazon.com/Antminer-S9-0-10W-Bitcoin-Miner/dp/B01GFEOV00](http://www.amazon.com/Antminer-S9-0-10W-Bitcoin-Miner/dp/B01GFEOV00)

**Using eclipse attacks.** If an eclipse attack is possible, the adversary can “pause” the auditor at a block height representing some previous state of the log, and can prevent the auditor from hearing about new blocks past this height. It is then free to mine blocks at its own pace, and can thus launch a successful split-view attack solely by mining  $k$  blocks, where  $k$  is the number of blocks the auditor requires to be mined after a block containing a given commitment in order to consider that commitment as valid. (It is standard in most Bitcoin wallets to use  $k = 6$ .)

Using our rough estimates above, it would cost the adversary 4,986 USD in electricity costs to mine a block, or 29,916 USD for  $k = 6$ . The hardware costs depends on how much time the adversary needs to conduct the attack, or how long they are able to continue their man-in-the-middle attack on the auditor. If—as a conservative number—the adversary wants to conduct the attack within a week, it must mine a block every 1.4 days to produce 6 blocks, which requires 1,285 mining rigs at a hardware cost of 3,084,000 USD. This brings the total cost of the attack to 3.1M USD, which is likely to deter at least a large fraction of potential adversaries. It may be affordable for a powerful adversary, however, but the attack is also fundamentally targeted: if the adversary wants to later compromise previously non-eclipsed auditors, it must mine a new set of blocks (assuming these auditors have more up-to-date blocks) and pay the electricity costs again. Even for an adversary with few financial constraints, this makes it significantly more difficult to conduct such an attack on a wide scale.

Furthermore, if the adversary takes 1.4 days to mine a block, or in general the auditor sees no new blocks until long after the expected 10-minute interval, it may assume that an eclipse attack is being performed. We can thus greatly increase the cost of the attack by adding simple checks to the auditor to ensure that there is a maximum interval between blocks. If we generously set such a check to require a maximum of 6 hours between blocks, then a total of 35,977 mining rigs are required at a cost of 86.3M USD.

In addition, the blocks must still follow the same difficulty level as honest blocks, so by mining these only in the eclipsed view of the network the adversary is not only expending the energy needed to do so but is also forfeiting the mining reward associated with them. As of February 2017, the Bitcoin mining reward is 12.5 bitcoins, or roughly 12,500 USD, so for  $k = 6$  the adversary must additionally forfeit 75,000 USD.

**Ignoring eclipse attacks.** If, for whatever reason, an eclipse attack is not possible, then an adversary can perform a split-view attack only if it can fork the Bitcoin blockchain. This naïvely requires it to control 51% of

the network’s mining power.

As of February 16 2017, the total hashing power of the Bitcoin network was 2,917,084 terahashes per second.<sup>8</sup> Conducting a 51% attack would therefore require the adversary to be able to compute more than  $\frac{2917084}{2}$  terahashes per second. Per hour, the total electricity cost would be  $\frac{2917084 \cdot 10^{12}}{2} \cdot 3600 \cdot J \cdot E$ , or—using our earlier estimates for  $J$  and  $E$ —14,437 USD per hour. In terms of hardware costs, if we use the figures for the Antminer S9 from before, the total number of mining rigs required would be greater than  $\frac{2917084 \cdot 10^{12}}{2 \cdot 14 \cdot 10^{12}} = 104186$ , at a total cost of at least 250M USD.

While more sophisticated attacks, such as selfish mining [10], have proposed strategies that fork the blockchain (also known as a “double spend” attack) using only 25% of the mining power, this would still require an investment of hundreds of millions of dollars. Such an attack would furthermore be highly visible, as the blockchain is regularly monitored for forks.

### Availability (S2)

While the decentralized (and thus fully replicated) nature of the blockchain can guarantee availability, it guarantees these properties only with respect to the commitments to statements made by the authority, rather than with respect to the statements—and thus the binaries—themselves. As discussed in Section 5.5, the use of the blockchain thus does not guarantee that binaries are actually available for inspection, or will continue to be into the future.

Using the archival nodes introduced in Section 5.5, we can achieve availability as long as these nodes are honest about whether or not they have mirrored the relevant data. Even if they do lie, monitors can still detect that an authority committed a statement without making the statement data available.

In the binary transparency setting, many ISPs and hosting providers already provide their customers local mirrors of Debian repositories. We therefore envision that ISPs can act as archival nodes on behalf of their hosting clients, which creates a decentralized network of archival nodes. We elaborate on the overheads required to do so in Sections 7.2 and 8.

### Auditor privacy (S3)

Recall from Section 4.2 that one of the goals of Con-tour was to avoid prolonged interactions and engage only in the atomic exchange of messages. In particular, the auditor receives pre-formed proofs of inclusion from the auditor (as opposed to requesting them for specific statements), retrieves commitments directly from the blockchain, does not engage in any form of gossip with monitors, and receives the latest block hash from

<sup>8</sup>[blockchain.info/charts/hash-rate](http://blockchain.info/charts/hash-rate)



Operation	Time complexity
commit	$O(n_S)$
prove_incl (one-time)	$O(\log(n_T))$
prove_incl (per statement)	$O(\log(n_S))$
check_incl	$O(\log(n_S) + \log(n_T))$

**Table 1:** Asymptotic computational costs for the operations of Contour, where  $n_S$  is the number of statements in a batch and  $n_T$  is the number of transactions in a block.

Object	Size Complexity
Inclusion proof	$O(\log(n_S) + \log(n_T))$
Log commitment (tx)	$O(1)$
Archival node data overhead	$O(n_S)$

**Table 2:** Asymptotic storage costs for the objects in Contour, where  $n_S$  is the number of statements in a batch and  $n_T$  is the number of transactions in a block.

archival nodes without providing any input of its own. We thus achieve privacy by design, as at no point in the process does the auditor reveal the statements in which it is interested to any other party.

One particular point to highlight is that Contour achieves auditor privacy despite the fact that auditors run SPV clients, which are known to potentially introduce privacy issues due to the use of Bloom filtering and the reliance on full nodes. This is because the proofs of inclusion contain both the raw transaction data and the block header, so the auditor does not need query a full node for the inclusion of the transaction and can instead verify it itself (and, as a bonus, saves the bandwidth costs of doing so).

## Deployability goals

### Efficiency (D1)

Table 1 summarizes the computational complexity of each of the operations required to run Contour, and Table 2 summarizes the size complexity.

As we will see in Sections 7.2 and 8, in real deployments of Contour there are already significant storage costs for the authority and archival nodes, as they must store the full set of binaries. It therefore does not impose a significant additional burden to have them perform relatively inefficient (i.e., linear in  $n$ ) operations or store relatively inefficient objects. As for the end-user devices on which the auditor is run, we impose a relatively minimal performance overhead (with everything logarithmic in  $n_S$  and/or  $n_T$ ), and confirm this in Section 7.2.3.

### Minimal setup (D2)

In terms of coordination, the only setup requirement in Contour is the role of the archival nodes, as the rest is just a matter of adding software. As we will see in Section 8 when we look at Debian, in some settings there are already natural candidates for these actors. More importantly, there are no trust requirements placed on these nodes to prevent log equivocation: even if archival nodes misbehave, monitors can still individually detect misbehavior by an authority that publishes commitments but not the underlying data.

### Comparison with existing solutions

To fully pinpoint both the benefits and tradeoffs of Contour, we compare it with several known systems designed to provide transparency. In particular, we consider the tradeoffs as compared to Certificate Transparency (CT), Collective Signing (CoSi) [27], CONIKS [21], and Bitcoin. We summarize these tradeoffs in Table 3.

Looking at Table 3, we first mention that the efficiency numbers for CoSi are somewhat misleading, as there is no global log and thus no notion of checking inclusion in the log; this is why the efficiency costs are constant. In fact, only Bitcoin and Contour ensure a globally consistent ledger, as certificates are stored in a distributed set of logs in CT and CONIKS and there is no proposed method for achieving consensus amongst them.

Arguably the main benefit of both CT and CONIKS is their efficiency, as the auditor is required to store only a single hash. The tradeoff, however, is that they cannot prevent the authority from launching a split-view attack, but instead rely on gossiping mechanisms to detect such misbehavior after the fact. As discussed in the introduction, this is problematic in a setting in which adversaries can launch persistent man-in-the-middle attacks. These systems also do not achieve robust privacy for the auditor, as it must periodically reveal information to the authority (or the monitor) about the objects in which it is interested.

The other main tradeoff we observe is, perhaps unsurprisingly, between efficiency and setup costs. The first three systems all require the establishment of some initial set of distributed entities—in the case of CT, log servers are essentially authorized by Google, in the case of CONIKS, identity providers are chosen by users and listed in a PKI, and in the case of CoSi, witnesses must form a Sybil-free set—that are trusted to some extent (if not individually, then as a group).

In contrast, in both Bitcoin and Contour, the blockchain is maintained by a decentralized network and is not subject to intervention by central authorities. While Contour mitigates the inefficiency of Bitcoin, it still requires the auditor to store some information from

	Security goals (S1-S3)			Deployability goals (D1-D2)		
	Split views	Availability	Auditor privacy	Efficiency (cost)	Efficiency (size)	Minimal setup
CT	detect	no*	no	$\log(n)$	1	no
CoSi	prevent	yes*	yes	1	1	no
CONIKS	detect	no	no	$\log(n)$	1	no
Bitcoin	prevent	yes	yes	$n$	$n$	yes
Contour	prevent	yes	yes	$\log(n)$	$b$	yes

**Table 3:** A comparison between existing solutions and Contour in terms of the five goals presented in Section 4.3. For efficiency, we measure the asymptotic costs for the auditor in terms of both the computations it must perform (‘cost’) and the data it must store (‘size’). We use  $n$  to denote the number of statements and  $b$  to denote the number (but not size) of blocks in the Bitcoin blockchain. For CoSi, availability is not a explicit requirement, but can be satisfied as long as at least one witness retains the data, and for CT it is not satisfied by the basic design but could be if auditors and monitors gossiped about certificates.

all the block headers. We demonstrate in the next section that Contour is nevertheless efficient enough to be practical, but leave it as an interesting open avenue of research to investigate to extent to which these tradeoffs between efficiency and decentralization are inherent.

## 7 Implementation and Performance

To test Contour and analyze its performance, we have implemented and provided benchmarks for a prototype Python module and toolset that developers can use. We have released the implementation as an open-source project, and will release details upon publication.

### Implementation details

The implementation consists of roughly 1000 lines of Python code, and provides a set of developer APIs and corresponding command-line tools. We used SHA-256 as the hashing algorithm to build Merkle trees, and modified versions of an existing Merkle tree implementation (<https://github.com/jvsteiner/merkletree>) and a Python-based Bitcoin library `pycoinnet` (<https://github.com/richardkiss/pycoinnet/>) in order to develop our Merkle tree and SPV client, respectively. (The Merkle tree modifications were necessary because the Merkle tree implementation in Bitcoin has a documented bug that must, for consensus reasons, be replicated in software using the Bitcoin protocol. The `pycoinnet` modifications were necessary as this library provides only the code to communicate using the Bitcoin protocol, rather than connecting to Bitcoin itself.)

**Authority:** We provide API calls for `Authority.commit`, which commits batches of statements to the Bitcoin blockchain, and `Authority.prove_incl`, which al-

lows it to generate inclusion proofs for individual statements.

**Auditor:** We provide an API call for `Auditor.check_incl`, which allows end-user software to verify proofs of inclusion. We also provide an `Auditor.sync` call that uses the Bitcoin SPV protocol to download and verify all the block hashes in the Bitcoin blockchain, so that inclusion proofs can be efficiently verified independently of third parties. (This call needs to be run only once.)

**Monitor:** We provide an API call for `Monitor.get_commits`, which gets all the statement batches associated with a specific authority. Monitors can then use these commitments to check the validity of the statement data (which they can retrieve from the authority or an archival node using a web server), and do whatever manual inspection is necessary; we consider this functionality outside of the scope of this paper.

**Archival node:** The archival node API can be used to operate an archival node, by specifying the authority’s Bitcoin address and web address where statement data is published. The archival state and mirrored statement data is stored as flat files on disk, allowing the archival node to provide access to auditors and monitors by running a web server. By accessing the archival state via a HTTPS server, auditors can securely authenticate the state of the archival node using public-key cryptography.

### Performance

To evaluate the performance of our implementation, we tested all the operations listed above on a laptop with an Intel Core i5 2.60GHz CPU and 12GB of RAM, that was connected to a WiFi network with an Internet connection of 5Mbit/s. We also assume that a batch to be commit-

Operation	Time ( $\mu$ s)	$\sigma$ ( $\mu$ s)
commit	<b>5.93 (s)</b>	<b>0.297 (s)</b>
prove_incl (one-time)	8.5	5.4
prove_incl (per statement)	12	6.4
check_incl	224	62.14

**Table 4:** Average time of individual operations, and standard deviation  $\sigma$ , where  $n$  is the number of statements in a batch. The timings for `commit` were averaged over 20 runs, and for `prove_incl` and `check_incl` over 1M runs. The timings for `commit` are in bold to emphasize that they are in seconds, not microseconds.

ted contains 1 million statements, although as was seen in Table 1 — and confirmed later on in Figure 4 — these numbers scale as expected (either logarithmically or linearly), so it is easy to extrapolate the results for other batch sizes given the ones we present here.

We consider the complexity of these operations in terms of their computational, storage, and bandwidth requirements. A summary of our timing benchmarks can be found in Table 4.

### Number of transactions per block

The overhead of both generating and verifying a proof of inclusion is dependent on the number of transactions in a Bitcoin block. To capture the worst-case scenario, we consider the maximum number of transactions that can fit into a block. Currently, the Bitcoin block size limit is 1 MB, up to 97 bytes of which is non-transaction data.<sup>9</sup> The minimum transaction size is 166 bytes,<sup>10</sup> so the upper bound on the number of transactions in a given block is 6,023. While this is far higher than the number of transactions that Bitcoin blocks currently contain,<sup>11</sup> we nevertheless use it as a worst-case cost and an acknowledgment that Bitcoin is evolving and blocks may grow in the future.

### Authority overheads

To run `commit` and `prove_incl`, an authority must have access to the full blocks in the Bitcoin blockchain, as well as the ability to broadcast transactions to the network. Rather than achieve these by running the authority as a full node, our implementation uses external blockchain APIs supplied by `blockchain.info` and `blockcypher.com`. This decision was based on the improved efficiency and ease of development for prototyping, but it does not affect the security of the system: authorities do not need to validate the blockchain, as in-

<sup>9</sup><https://en.bitcoin.it/wiki/Block>

<sup>10</sup>[https://en.bitcoin.it/wiki/Maximum\\_transaction\\_rate](https://en.bitcoin.it/wiki/Maximum_transaction_rate)

<sup>11</sup><https://blockchain.info/charts/n-transactions-per-block>

valid blocks from a dishonest external API simply result in invalid inclusion proofs that are rejected by the auditor.

To run `commit`, an authority must first build the Merkle tree containing its statements. Sampled over 20 runs, the average time to build a Merkle tree for 1M statements was 5.9 s ( $\sigma = 0.29$  s). After building the tree, an authority next embeds its root hash (which is 32 bytes) into an `OP_RETURN` Bitcoin transaction to broadcast to the network. Sampled over 1,000 runs, the average time to generate this transaction — in the standard case of one input and two outputs, one for `OP_RETURN` and one for the authority’s change — was 0.03 s ( $\sigma = 0.007$  s). The average total time to run `commit` was thus 5.93 s, as seen in Table 4, and it resulted in 235 bytes (the size of the transaction) being broadcast to the network.


Next, to run `prove_incl`, the authority proceeds in two phases: first constructing the Merkle proof for their transaction within the block where it eventually appears, and next constructing the Merkle proof for each statement represented in a transaction. The time for the first phase, averaged over 1M runs and for a block with 6,023 transactions (our upper bound from Section 7.2.1), was 8.5  $\mu$ s. This is denoted “one-time” in Table 4 as it is done only once per batch. The time for the second phase, averaged over 1M runs, was 12  $\mu$ s for each individual statement (thus denoted “per statement” in Table 4). Generating inclusion proofs for all the statements in the batch would thus take around 12 s. In terms of bandwidth and storage, the block may be up to 1 MB in size. In terms of the memory costs, the size of the Merkle tree for 1M leafs in memory is 649MB.

Additionally, in order to ensure that its transaction makes it into a block quickly, the authority may want to pay a fee. The current recommended rate is 120 satoshis/byte (<https://bitcoinfees.21.co/>), so for a 235-byte transaction the authority can expect to pay 28,200 satoshis. As of February 2017, this is roughly 0.28 USD.

### Auditor overheads

For the auditor, we considered two costs: the initial cost to retrieve the necessary header data (`sync`), and the cost to verify an inclusion proof (`check_incl`). We do not provide performance benchmarks for the `Auditor.get_arch_state` call, as this is a simple web request that returns a single 32-byte hash.

To run `sync`, auditors use the Bitcoin SPV protocol to download and verify the headers of each block, which are 80 bytes each. As of this writing on February 14 2017, there are 453,021 valid mined blocks, which equates to 36.2 MB of block headers. Once downloaded, however, the auditor needs to keep only the 32-byte block hash, so only 14.5 MB of data needs to be stored on disk. Going



**Figure 4:** The time to verify an inclusion proof with varying numbers of statements in the batch, averaged over 100K runs.

forward, the Bitcoin network generates approximately 144 blocks per day, so the amount of downloaded data will increase by 11.5 kB daily, and the amount of stored data by 4.6 kB daily.

To verify the validity of the block headers in the chain, the client must perform one SHA-256 hash per block header; on average, it took us on average 116 seconds (over 5 runs) for the Python SPV client to download and verify all the block headers from the network. This initial bootstrapping process needs to be performed only once per auditor.

To run `check_incl`, we again use our upper bound from Section 7.2.1 and assume every block contains 6,023 transactions. This means the inclusion proof contains: (1) an 80-byte block header; (2) the raw transaction data, which is 235 bytes; (3) a Merkle proof for the transaction, which consists of  $\log(6,023) - 1$  32-byte hashes (the root hash is already provided in the block header); and (4) a Merkle proof for the statement, which consists of  $\log(1,000,000) - 2$  32-byte hashes (the root hash is already provided in the transaction data, and the auditor computes the statement hash itself). The total bandwidth cost is therefore around 1275 bytes. Averaged over 1M runs, the time for the auditor to verify the inclusion proof was 224  $\mu$ s ( $\sigma = 62.14 \mu$ s).

To confirm that the time to run `check_incl` scales logarithmically with the number of statements in the batch, we also changed this number. The results are in Figure 4 (and do confirm logarithmic scaling).

### Monitor overheads

Monitors must run a Bitcoin full node in order to get a complete uncensored view of the blockchain. As of Jan-

uary 2017, running a Bitcoin full node requires 120 GB of free disk space, increasing by up to 144 MB daily. It took us around three days to fully bootstrap a Bitcoin full node and verify all the blocks, although again this operation needs to be performed only once per monitor.

### Archival nodes overheads

Like monitors, archival nodes need to run a Bitcoin full node. Additionally, archival nodes must download and store all the statement data from the authority. The costs here are entirely dependent on the number and size of the statements; we examine the costs for Debian in Section 8.

As in Debian, many applications use statements that represent files. We may therefore expect that, in addition to a Merkle tree, authorities would use metadata files to link each leaf in the tree to a file on the server that archival nodes then mirror; this would be particularly useful in a setting — like Debian — where it would be undesirable to reorganize files that are already stored. The metadata file would consist of a mapping of 32-byte hashes to filenames. Assuming filenames are no longer than 64 bytes each, this would introduce an extra storage overhead, for both authorities and archival nodes, of up to 96 bytes per statement.

## 8 Use Case: Debian

To go beyond basic benchmarks and analyze the operation of Contour on a real system, we used it to audit software binaries in the Debian repository. Our results show that, as desired, Contour provides a way to add transparency to this repository without major changes to the existing infrastructure and with minimal overheads.

We extracted the software package metadata for all processor architectures and releases of Debian from the Debian FTP archive (<https://www.debian.org/mirror/ftpmirror>) over a one-week period from January 20-27 2017. The archive is updated four times a day. At the beginning of this period there were 976,214 unique software binaries available for download from the Debian software repositories, constituting 1.7 TB of data, and by the end there were 980,469.

To initiate the system, we first committed to all the existing 976,214 software packages. The Debian package metadata already contains the SHA-256 hashes for these packages, so we only needed to build a Merkle tree from these hashes (rather than compute them ourselves first). This took approximately 6 seconds (which is in line with our benchmarks in Table 4 for 1M statements).

As the archive was updated, we kept track of the package hashes being added and created a new batch for each update. The average batch size was 1,040, and the average time to build a Merkle tree for the batch was 0.0052 seconds.

Recall from Section 7.2.2 that committing one transaction to the blockchain costs roughly 0.28 USD in fees, so this would cost 1.12 USD per day.

In terms of overhead for archival nodes, to fully satisfy availability they must store all the data from the authority, as well as deleted packages (which should be monitored as well). This means storing 1.7 TB, and an additional average of 11 GB per day, or 4 TB per year. There are already 269 Debian mirrors hosting the full 1.7 TB data set,<sup>12</sup> and at least one mirror hosting all the deleted packages too, effectively acting as an archival node.<sup>13</sup> This is by far the highest overhead incurred by our system, and we expect that only a small number of mirrors would have the storage capacity to run an archival node.

As discussed in Section 7.2.5, we can also enable archival nodes to rebuild Merkle trees with minimal changes to the existing Debian archive infrastructure. This requires storing only an additional 84 kB metadata file per batch (containing the mapping from hashes to filenames), and an initial 79 MB metadata file. These metadata files consist of a mapping of hashes of software packages to their filenames in the Debian archives.

Finally, the proof of inclusion of each software package would need to be added to the software package (.deb) files as metadata when downloaded by a Debian device using a command such as `apt install`. At 980K software packages, this would require a maximum of 1.3 kB of extra storage and bandwidth for end-user devices per package downloaded. Given that the average package size is 1337 kB, this is only a 0.1% overhead. 1.3 GB of extra storage is required for Debian repository mirrors to store the proofs of inclusion, which is only a 0.07% overhead.

## 9 Discussion and Extensions

**IPFS.** The InterPlanetary File System (IPFS)<sup>14</sup> is a peer-to-peer data distribution protocol, where content can be addressed by its hash. In theory, this would be a useful addition to Contour, as auditors and monitors would be able to determine the IPFS address of binary data simply from the Merkle root stored in the blockchain, thus eliminating the need to manually seek out archival nodes. The maximum size of an IPFS block is 1 MB, however, so while IPFS may be directly useful in the context of certificate (or key) transparency, it loses its usefulness in the context of binary transparency because binaries would have to be split into multiple IPFS blocks, thus making it impossible to address binaries purely based on their hash.

<sup>12</sup>[www.debian.org/mirror/list](http://www.debian.org/mirror/list)

<sup>13</sup>[snapshot.debian.org/](http://snapshot.debian.org/)

<sup>14</sup>[ipfs.io/](http://ipfs.io/)

**Selective disclosure.** When releasing software updates that patch critical security vulnerabilities, some software vendors may prefer not to reveal to potential attackers that, in the window of time in which a commitment has not yet been included in the blockchain, they can take advantage of victims with this vulnerable software installed. In such a case, Contour accounts for this by allowing the authority to commit to a batch of binaries visibly on the blockchain, but delay the publication of the binaries themselves until the commitment is sufficiently deep in the blockchain.

**Generalized transparency.** Although we have designed Contour for the specific application of binary transparency, the system is general enough to be applied to other applications requiring transparency. With the tradeoffs discussed in Section 6.3, it can even be applied to the setting of certificate transparency, using CAs as authorities.

**Archival node scalability.** The current design of Contour requires archival nodes to store all data, which as we have seen incurs a significant overhead. There are likely many alternative designs that alleviate these requirements, such as a *sharded* solution in which archival nodes store only the data for which they sufficient space, and we leave an exploration of this space as an interesting open problem.

## 10 Conclusion

We have proposed Contour, a system that provides proactive transparency, scales logarithmically for auditors, and does not require the initial coordination of a Sybil-free set of nodes. To the best of our knowledge, it is also the first system for providing binary transparency.

We have demonstrated that, even for attackers that are capable of performing (for free) persistent man-in-the-middle attacks and are targeted a single device, compromising the integrity of the system requires roughly 85M USD in energy and hardware costs. We also saw that Contour could be applied today to the Debian software repository with relatively minimal changes and overhead to existing infrastructure, with the main extra cost being the storage requirements of archival nodes that mirror the repository data. The overheads for end users, in contrast, are quite minimal, with the proof of inclusion for a binary within a batch of size 1M being only 1.3 kB and taking only 224  $\mu$ s to verify.

## References

- [1] M. Apostolaki, A. Zohar, and L. Vanbever. Hijacking Bitcoin: Large-scale Network Attacks on Cryptocurrencies, 2016. [arxiv.org/abs/1605.07524](https://arxiv.org/abs/1605.07524).

- [2] M. Bartoletti and L. Pompianu. An analysis of Bitcoin OP\_RETURN metadata. In *4th Workshop on Bitcoin and Blockchain Research*, 2017. To appear.
- [3] D. Basin, C. Cremers, T. H.-J. Kim, A. Perrig, R. Sasse, and P. Szalachowski. ARPKI: Attack Resilient Public-Key Infrastructure. In *Proceedings of ACM CCS 2014*, pages 382–393, 2014.
- [4] J. Bonneau. EthIKS: Using Ethereum to audit a CONIKS key transparency log. In *3rd Workshop on Bitcoin and Blockchain Research*, 2016.
- [5] J. Bonneau, A. Miller, J. Clark, A. Narayanan, J. A. Kroll, and E. W. Felten. Research perspectives and challenges for Bitcoin and cryptocurrencies. In *Proceedings of the IEEE Symposium on Security and Privacy*, 2015.
- [6] M. Chase and S. Meiklejohn. Transparency Overlays and Applications. In *ACM SIGSAC Conference on Computer and Communications Security*, 2016.
- [7] L. Chuat, P. Szalachowski, A. Perrig, B. Laurie, and E. Messeri. Efficient Gossip Protocols for Verifying the Consistency of Certificate Logs. In *IEEE Conference on Communications and Network Security*, 2015.
- [8] B. Dowling, F. Günther, U. Herath, and D. Stebila. Secure Logging Schemes and Certificate Transparency. In *Proceedings of ESORICS 2016*, 2016.
- [9] A. Eijdenberg, B. Laurie, and A. Cutter. Verifiable Data Structures, 2015. [github.com/google/trillian/blob/master/docs/VerifiableDataStructures.pdf](https://github.com/google/trillian/blob/master/docs/VerifiableDataStructures.pdf).
- [10] I. Eyal and E. G. Sirer. Majority is not enough: Bitcoin mining is vulnerable. In *Financial Cryptography and Data Security*, 2014.
- [11] C. Farivar. Judge: Apple must help FBI unlock San Bernardino shooter’s iPhone, 2016. [arstechnica.com/tech-policy/2016/02/judge-apple-must-help-fbi-unlock-san-bernardino-shooters-iphone/](http://arstechnica.com/tech-policy/2016/02/judge-apple-must-help-fbi-unlock-san-bernardino-shooters-iphone/).
- [12] C. Fromknecht, D. Velicanu, and S. Yakubov. A decentralized public key infrastructure with identity retention. IACR Cryptology ePrint Archive, Report 2014/803, 2014. <http://eprint.iacr.org/2014/803.pdf>.
- [13] A. Gervais, H. Ritzdorf, G. Karame, and S. Capkun. Tampering with the Delivery of Blocks and Transactions in Bitcoin. In *Proceedings of ACM CCS 2015*, 2015.
- [14] D. Goodin. Google warns of unauthorized TLS certificates trusted by almost all OSes, 2015. [arstechnica.com/security/2015/03/google-warns-of-unauthorized-tls-certificates-trusted-by-almost-all-oses/](http://arstechnica.com/security/2015/03/google-warns-of-unauthorized-tls-certificates-trusted-by-almost-all-oses/).
- [15] E. Heilman, A. Kendler, A. Zohar, and S. Goldberg. Eclipse Attacks on Bitcoin’s Peer-to-Peer Network. In *Proceedings of USENIX Security 2015*, 2015.
- [16] T. H.-J. Kim, L.-S. Huang, A. Perrig, C. Jackson, and V. Gligor. Accountable key infrastructure (AKI): a proposal for a public-key validation infrastructure. In *Proceedings of WWW 2013*, pages 679–690, 2013.
- [17] E. K. Kogias, P. Jovanovic, N. Gailly, I. Khoffi, L. Gasser, and B. Ford. Enhancing Bitcoin Security and Performance with Strong Consistency via Collective Signing. In *Proceedings of USENIX Security 2016*, 2016.
- [18] B. Laurie, A. Langley, and E. Kasper. Certificate transparency, 2013.
- [19] J. Leyden. Inside ‘Operation Black Tulip’: DigiNotar hack analysed, 2011. [www.theregister.co.uk/2011/09/06/diginotar\\_audit\\_damning\\_fail/](http://www.theregister.co.uk/2011/09/06/diginotar_audit_damning_fail/).
- [20] S. Matsumoto and R. M. Reischuk. IKP: Turning a PKI Around with Blockchains. IACR Cryptology ePrint Archive, Report 2016/1018, 2016. [eprint.iacr.org/2016/1018](http://eprint.iacr.org/2016/1018).
- [21] M. S. Melara, A. Blankstein, J. Bonneau, E. W. Felten, and M. J. Freedman. CONIKS: Bringing Key Transparency to End Users. In *Proceedings of USENIX Security 2015*, 2015.
- [22] S. Nakamoto. Bitcoin: A Peer-to-Peer Electronic Cash System, 2008. <https://bitcoin.org/bitcoin.pdf>.
- [23] E. Nakashima, G. Miller, and J. Tate. U.S., Israel developed Flame computer virus to slow Iranian nuclear efforts, officials say, 2012. [www.washingtonpost.com/world/national-security/us-israel-developed-computer-virus-to-slow-iranian-nuclear-efforts-officials-say/2012/06/19/gJQA6xBPoV\\_story.html](http://www.washingtonpost.com/world/national-security/us-israel-developed-computer-virus-to-slow-iranian-nuclear-efforts-officials-say/2012/06/19/gJQA6xBPoV_story.html).
- [24] L. Nordberg, D. Gillmor, and T. Ritter. Gossiping in CT, 2016. [tools.ietf.org/html/draft-ietf-trans-gossip-03](https://tools.ietf.org/html/draft-ietf-trans-gossip-03).
- [25] M. D. Ryan. Enhanced Certificate Transparency and End-to-end Encrypted Mail. In *Proceedings of NDSS 2014*, 2014.
- [26] A. Singh, T.-W. J. Ngan, P. Druschel, and D. S. Wallach. Eclipse attacks on overlay networks: Threats and defenses. In *IEEE Conference on Computer Communications*, 2006.
- [27] E. Syta, I. Tamas, D. Visher, D. I. Wolinsky, P. Jovanovic, L. Gasser, N. Gailly, I. Khoffi, and B. Ford. Keeping Authorities “Honest or Bust” with Decentralized Witness Cosigning. In *IEEE Symposium on Security and Privacy (“Oakland”)*, 2016.
- [28] A. Tomescu and S. Devadas. Catena: Preventing Lies with Bitcoin. IACR Cryptology ePrint Archive, Report 2016/1062, 2016.